

456: Network-Centric Programming

Yingying (Jennifer) Chen

Web: <https://www.winlab.rutgers.edu/~yychen/>
Email: yingche@scarletmail.rutgers.edu

Department of Electrical and Computer Engineering
Rutgers University

File Pointers and Seeking

- The system remembers the position in a (real) file through a file pointer (32-bit offset), which automatically advances when you read or write
- These functions can query or modify the position of the pointer:

```
long ftell(FILE *fp);
int fseek(FILE *fp, long offset, int whence);
    // SEEK_SET, SEEK_CUR, SEEK_END
void rewind(FILE *fp);
```

fseek() Example

```
/* fseek example */
#include <stdio.h>
// This program modifies a specific part of a file using fseek()

int main () {
    FILE * pFile;
    pFile = fopen ( "example.txt" , "w" );
    fputs ( "This is an apple." , pFile );
    fseek ( pFile , 9 , SEEK_SET );
    fputs ( " sam" , pFile );
    fclose ( pFile );
    return 0;
}
```

After this code is successfully executed, the file example.txt contains:

This is a sample.

Low-level seeking example -- File with hole? code/lecture_4_code/file_hole

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include "ourhdr.h"

char buf1[] = "abcdefghijkl";
char buf2[] = "ABCDEFGHIJ";
#define FILE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH) // define file access permission

int
main(void)
{
    int fd;

    if ((fd = creat("file.hole", FILE_MODE)) < 0)
        err_sys("creat error");

    if (write(fd, buf1, 10) != 10)
        err_sys("buf1 write error");
    /* offset now = 10 */

    if (lseek(fd, 40, SEEK_SET) == -1)
        err_sys("lseek error");
    /* offset now = 40 */

    if (write(fd, buf2, 10) != 10)
        err_sys("buf2 write error");
    /* offset now = 50 */

    exit(0);
}
```

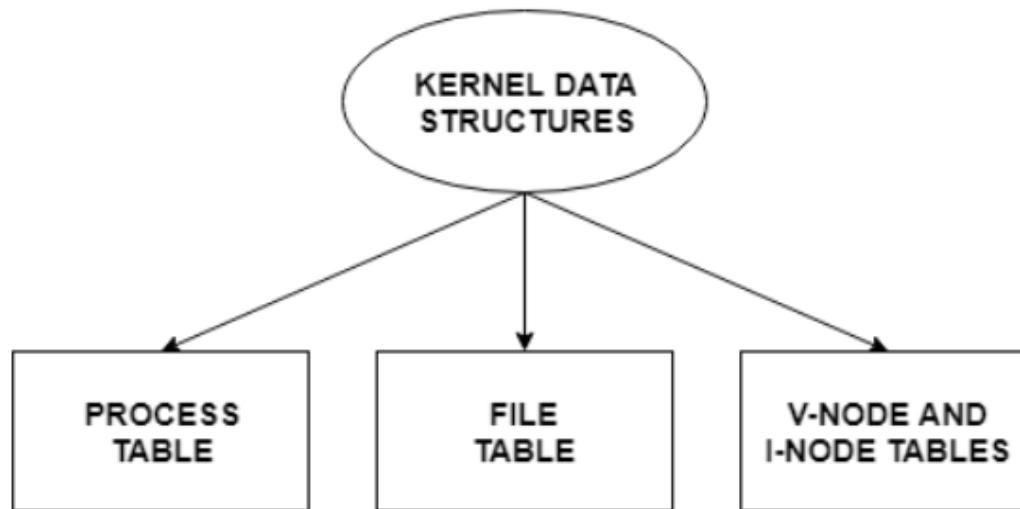
Split to 2

%od -c file.hole

-od command to look at the contents of the file.
-The -c flag tells it to print the contents as characters.

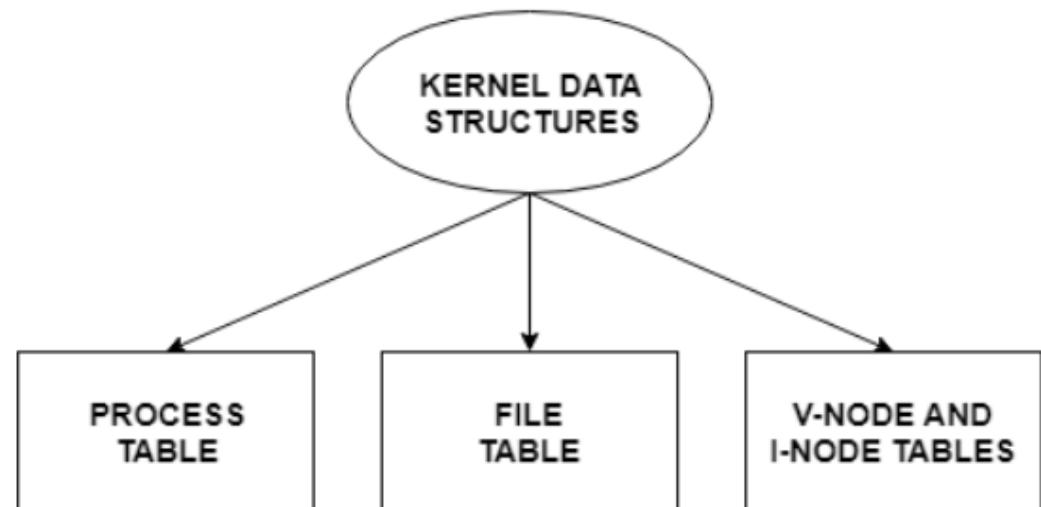
File Sharing

- Files can be opened multiple times
 - By same or different processes
- Kernel Data Structures



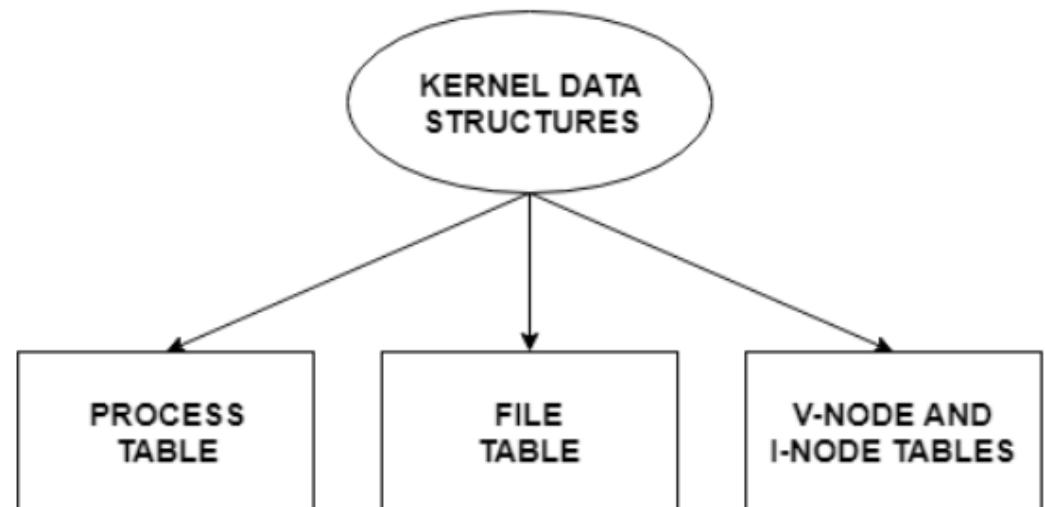
Kernel File Data Structures

- Process Table
 - Process table stores information about all the running processes, including the storage information, execution status, file information etc.
- File Table
 - File Table contains entries about all the files including file status, file offset

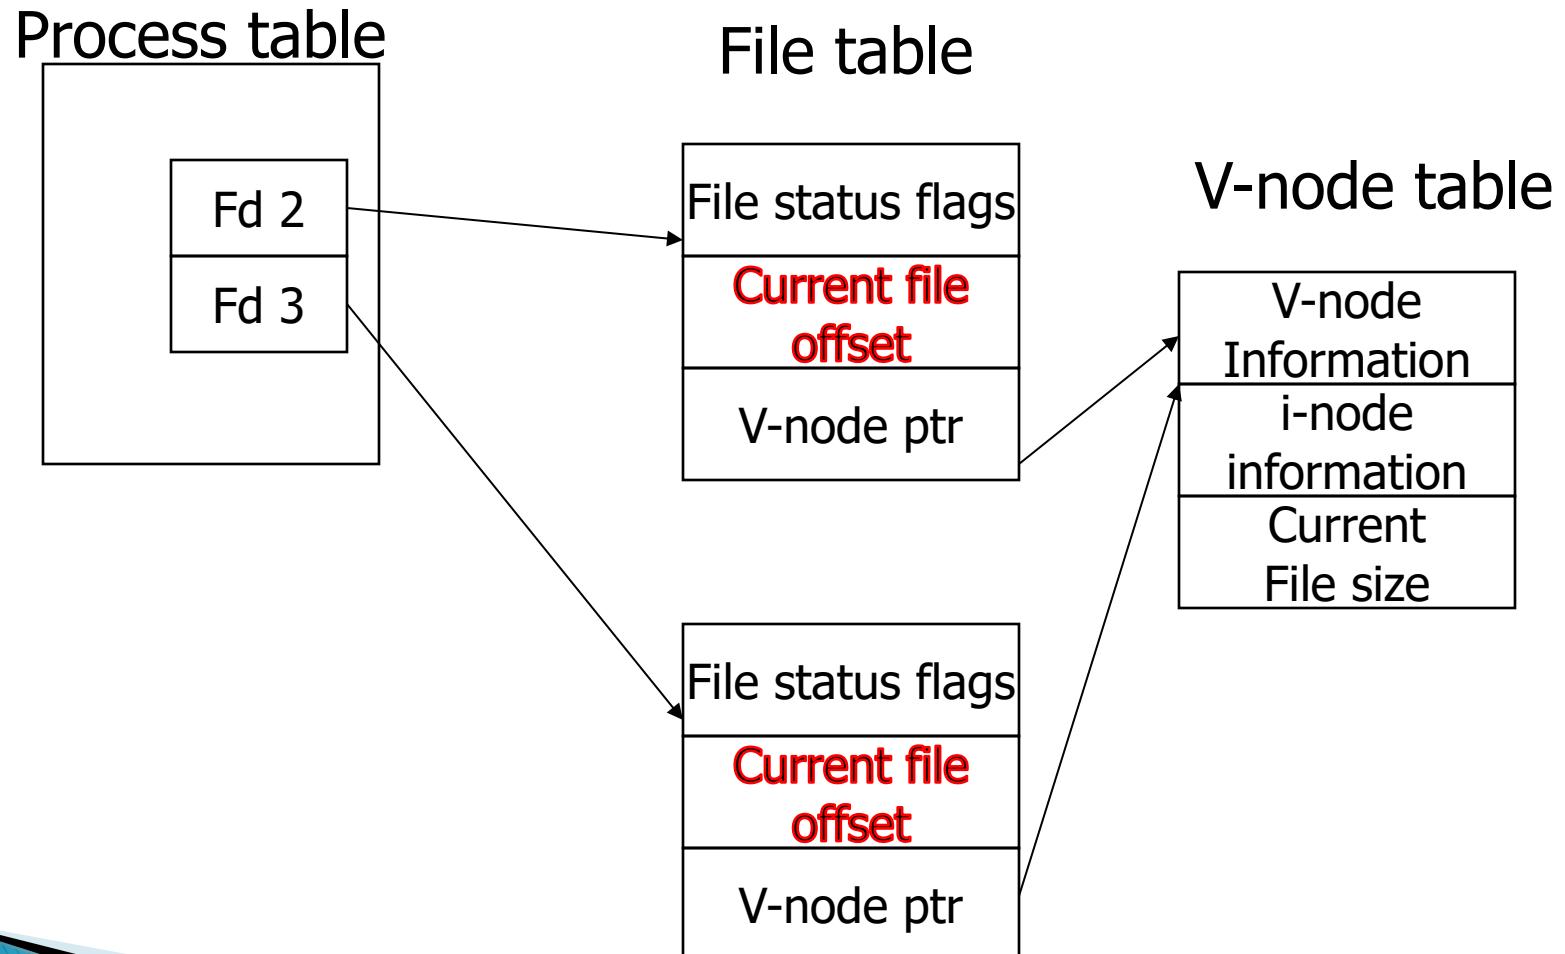


Kernel File Data Structures

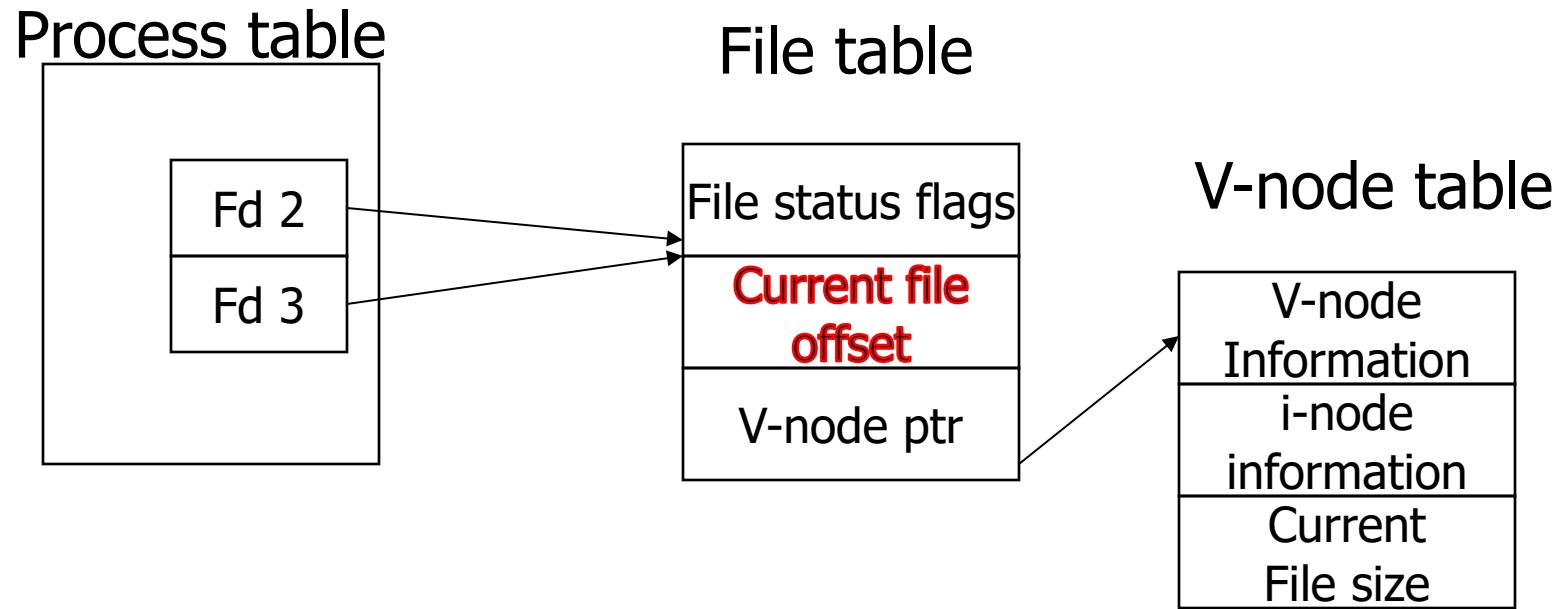
- v-node
 - The v-node is an abstract concept that defines the method to access file data without worrying about the actual structure of the system.
- i-node
 - The i-node specifies file access information like file storage device, read/write procedures etc.



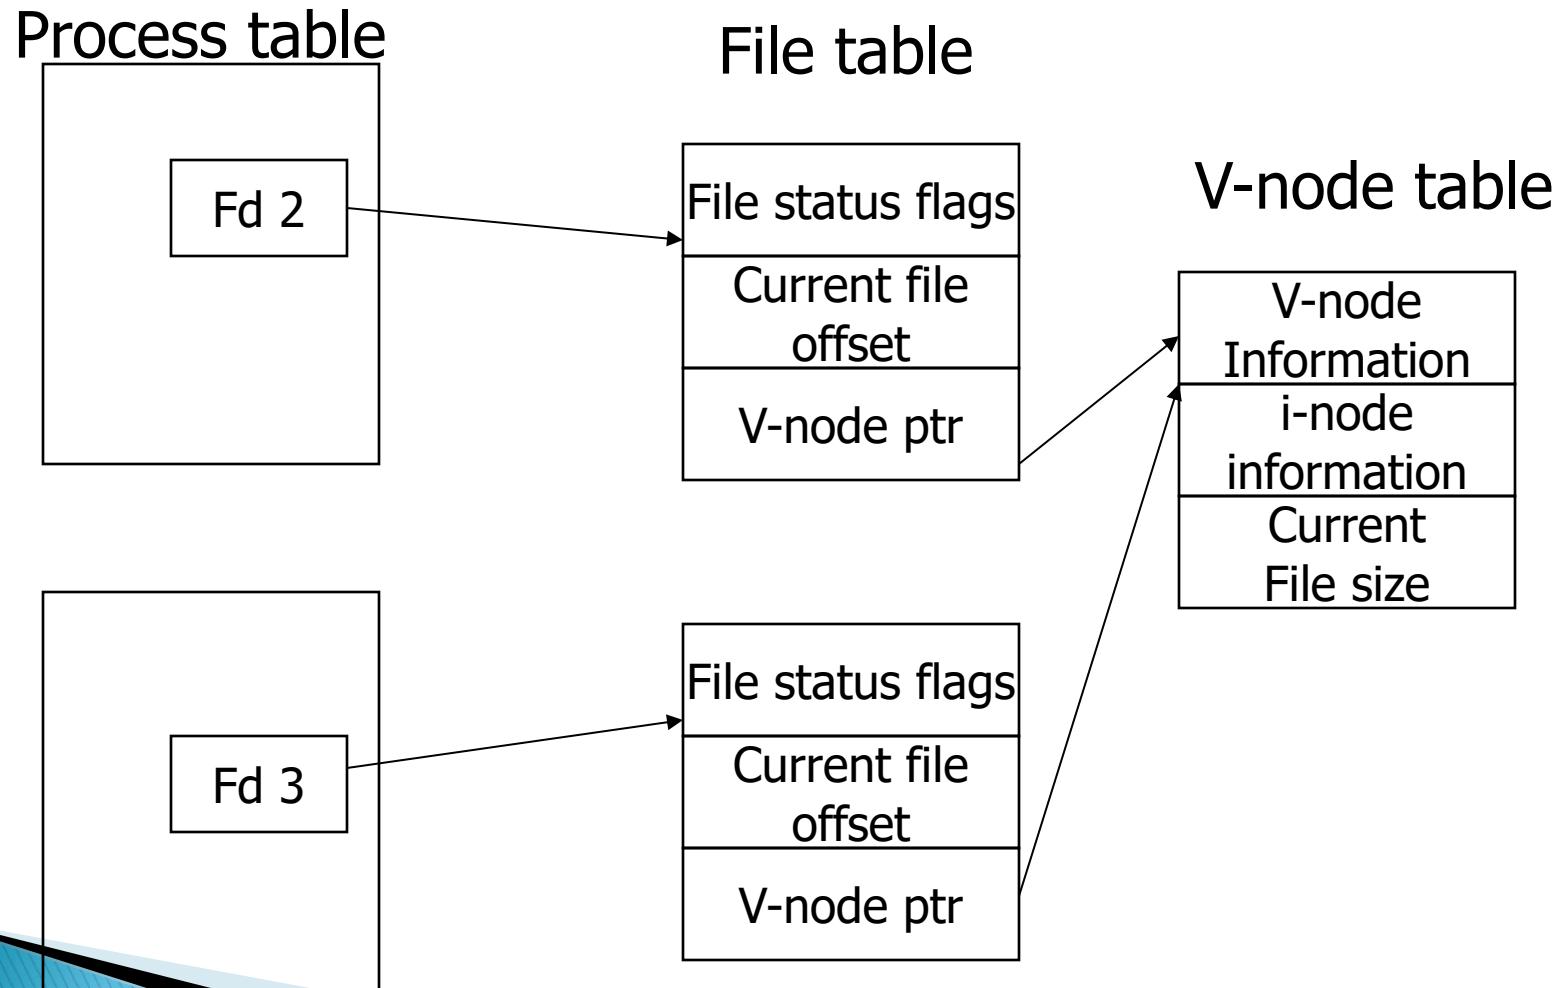
Kernel File Data Structures – A process opening the same file twice



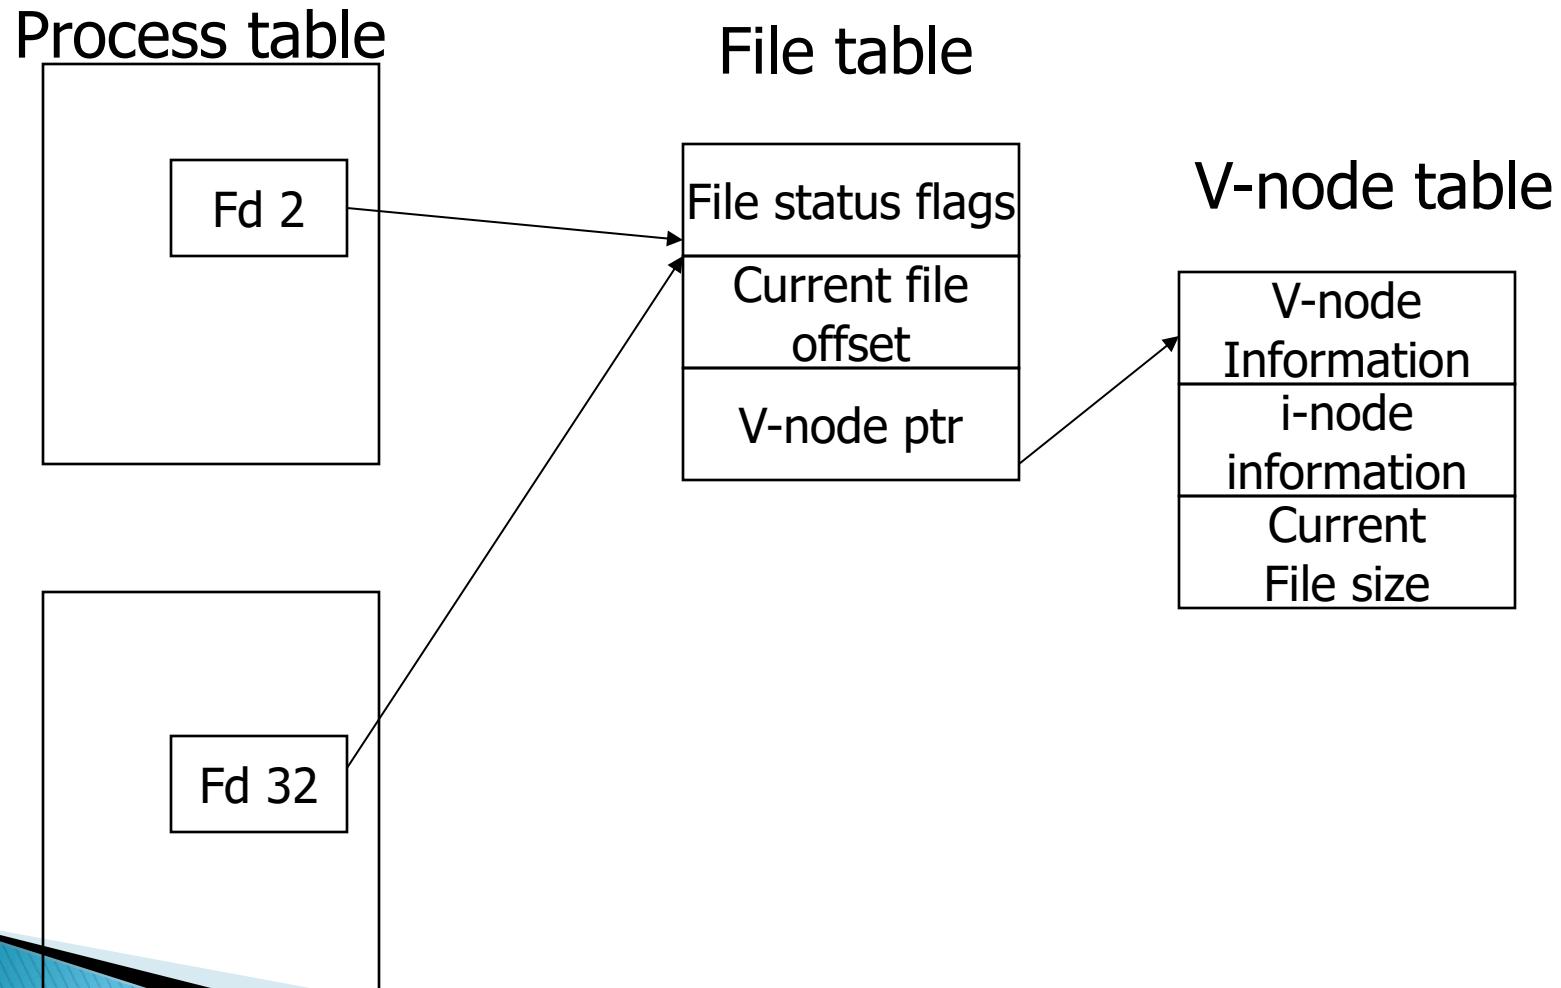
Kernel File Data Structures – A process with a copied file descriptor



Kernel File Data Structures – Two processes opening the same file



Kernel File Data Structures – A child process inheriting the same file



Directories and File Metadata

Getting file metadata

- ▶ Metadata is the “data about data”. It is used for summarizing basic information of data.
- ▶ `int stat(char* filename, struct stat* buf)`
 - Fills in the stat data structure with information about file type, size, permissions, time, ...

Struct stat

```
/* Metadata returned by the stat and fstat functions */
struct stat {
    dev_t          st_dev;        /* Device */
    ino_t          st_ino;        /* inode */
    mode_t         st_mode;       /* Protection and file type */
    nlink_t        st_nlink;      /* Number of hard links */
    uid_t          st_uid;        /* User ID of owner */
    gid_t          st_gid;        /* Group ID of owner */
    dev_t          st_rdev;       /* Device type (if inode device) */
    off_t          st_size;       /* Total size, in bytes */
    unsigned long  st_blksize;    /* Blocksize for filesystem I/O */
    unsigned long  st_blocks;     /* Number of blocks allocated */
    time_t         st_atime;      /* Time of last access */
    time_t         st_mtime;      /* Time of last modification */
    time_t         st_ctime;      /* Time of last change */
};
```

File types (mode)

- ▶ Regular file
- ▶ Directory file
- ▶ Character special file (device access, e.g. serial port)
- ▶ Block special file (device access, e.g., disk)
- ▶ FIFO (pipe)
- ▶ Socket (network connection)
- ▶ Symbolic link (pointer to another file)

File types (mode)

- ▶ Use MACROS to test
 - `S_ISREG(m)` is it a regular file?
 - `S_ISDIR(m)` directory?
 - `S_ISCHR(m)` character device?
 - `S_ISBLK(m)` block device?
 - `S_ISFIFO(m)` fifo?
 - `S_ISLNK(m)` symbolic link?
 - `S_ISSOCK(m)` socket?

Example

code/lecture_4_code/stat

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>

//Querying and manipulating a file's st_mode bits and file size.

int main (int argc, char **argv){
    struct stat st;
    char *type, *readok;
    stat(argv[1], &st); //extract metadata
    if (S_ISREG(st.st_mode)) /* Determine file type */
        type = "regular";
    else if (S_ISDIR(st.st_mode))
        type = "directory";
    else
        type = "other";
    if ((st.st_mode & S_IRUSR)) /* Check read access */
        readok = "yes";
    else
        readok = "no";                                % ./main data_1MB.txt
                                                % ./main data_10bytes.txt

    printf("type: %s, read: %s\n", type, readok);
    printf("File size: %lld\n", st.st_size);
    return 0;
}
```

**Why does the operating system need
to distinguish these file types?**

File types

- Unix recognizes a number of different file types.
- A *regular file* contains some sort of binary or text data.
- To the kernel there is no difference between text files and binary files.
- A *directory file* contains information about other files.
- A *socket* is a file that is used to communicate with another process across a network

Unix Filesystem Structure

Directory manipulation

- ▶ `mkdir, rmdir`
 - create and remove
- ▶ `opendir, readdir, rewinddir, closedir`
 - Read directory entries
- ▶ There is no direct write
 - Implicit in ‘create file’ and ‘`mkdir`’
- ▶ `chdir, getcwd`
 - Set and get working directory for current process

Directory structure

```
#include <dirent.h>

DIR *opendir(const char *pathname);
DIR *fdopendir(int fd);

struct dirent *readdir(DIR *dp);

void rewinddir(DIR *dp);
int closedir(DIR *dp);

long telldir(DIR *dp);

void seekdir(DIR *dp, long loc);
```

Both return: pointer if OK, NULL on error

Returns: pointer if OK, NULL at end of directory or error

Returns: 0 if OK, -1 on error

Returns: current location in directory associated with *dp*

The DIR structure is an internal structure used by these seven functions to maintain information about the directory being read.

- The purpose of the DIR structure is similar to that of the FILE structure maintained by the standard I/O library

Accessing system calls through library – ls

code/lecture_4_code/directory

```
#include <sys/types.h>
#include <dirent.h>
// This program read the file names in a directory.

int
main(int argc, char *argv[])
{
    DIR          *dp;
    struct dirent  *dirp;

    if (argc != 2)
        err_quit("a single argument (the directory name) is required");

    if ( (dp = opendir(argv[1])) == NULL)
        err_sys("can't open %s", argv[1]);                                % ./main dir

    while ( (dirp = readdir(dp)) != NULL)
        printf("%s\n", dirp->d_name);

    closedir(dp);
    exit(0);
}
```

Common file operations

- ▶ access – access control check
- ▶ chdir – change directory
- ▶ chown – change owner
- ▶ opendir – open directory
- ▶ remove – delete file or directory
 - Also unlink (file only), rmdir (directory only)
- ▶ rename – rename directory

access – determine accessibility of a file

```
#include <unistd.h>

int access(const char *path, int amode);
```

- ▶ *access()* function checks the file named by the pathname pointed to by the *path* argument for accessibility according to the bit pattern contained in *amode*.
- ▶ *amode*: (R_OK, W_OK, X_OK) or the existence test (F_OK).

The following example tests whether a file named **myfile** exists in the **/tmp** directory.

```
#include <unistd.h>
...
int result;
const char *filename = "/tmp/myfile";

result = access (filename, F_OK);
```

chdir, opendir

chdir – change working directory

```
#include <unistd.h>

int chdir(const char *path);
```

chdir() causes the directory named by the pathname pointed to by the *path* argument to become the current working directory;

opendir – open a directory

```
#include <dirent.h>
```

```
DIR *opendir(const char *dirname);
```

opendir() function opens a directory stream corresponding to the directory named by the *dirname* argument.

chown, remove

chown – change file owner and group

```
#include <unistd.h>
int chown(const char * pathname, uid_t owner, gid_t group);
```

`chown` changes a file's user ID and group ID, but if either of the arguments *owner* or *group* is -1 , the corresponding ID is left unchanged.

remove – remove a file/directory

```
#include <stdio.h>
```

```
int remove(const char * path);
```

► `remove()` causes the file named by the pathname pointed to by *path* to be deleted.

► If *path* name is a directory, `remove(path)` is equivalent to `rmdir(path)`.

rename

rename – rename a file

```
#include <stdio.h>

int rename(const char *old, const char *new);
```

rename() changes the name of a file. The *old* argument points to the pathname of the file to be renamed.
The *new* argument points to the new pathname of the file.

Example

code/lecture_4_code/chdir

```
#include <sys/types.h>
#include <dirent.h>

//Change directory to /tmp using chdir()

int main(void){
    if (chdir("/tmp") < 0)
        err_sys("chdir failed");
    printf("chdir to /tmp succeeded\n");
    return 0;
}
```

```
%pwd
% ./main
%pwd
```

The current working directory didn't change.

- Each program runs in a separate process, so the current working directory of the shell is unaffected by the call to *chdir* in the program.
- To change the directory from the shell, *cd* command needs to be invoked. *chdir* function is built inside of the *cd* command.

File System Links

Hard links

- Every directory entry points to an i-node (which represents a file)
 - Multiple entries can point to the same file (hard links)
- Linking to files is possible

Hard links

- ▶ *link(existingpath, newpath)* creates additional directory entry to existing file
- ▶ *unlink* removes the link

```
#include <unistd.h>

int link(const char *existingpath, const char *newpath);

int unlink(const char *pathname);
```

Symbolic links

- Represented by special files
- Also called Symlinks/Soft links
- Can span filesystems, users can create links to both directories and links

Symbolic links

- ▶ `symlink(actualpath, sympath)` – creates link
- ▶ `readlink(pathname, ...)` – reads `sympath`

```
#include <unistd.h>

int symlink(const char *actualpath, const char *sympath);
ssize_t readlink(const char* restrict pathname, char *restrict buf,
                 size_t bufsize);
```

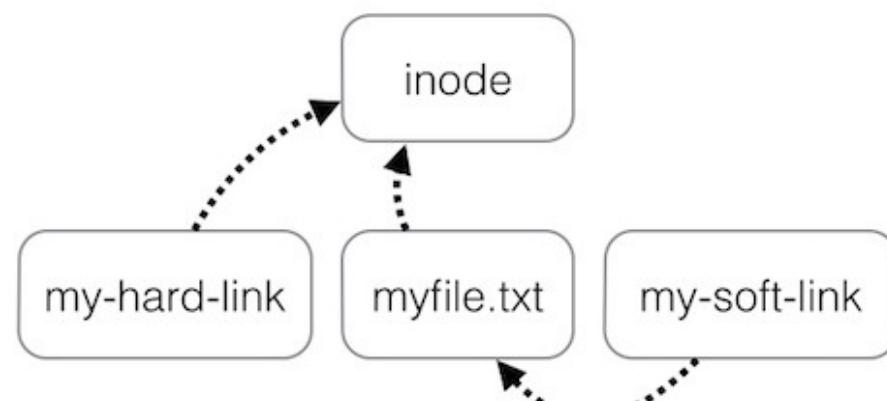
Two Types of File System Links

■ Hard Links

- A hard link is pointed directly to the *i-node* of the file.

■ Symbolic Links

- A symbolic link is an indirect pointer to a file



Example in Bash: Hard links

code/lecture_4_code/try

```
$ ln dir/data.txt mylink          create a hard link
$ cat mylink
1234567890                         show the content in data.txt
$ ls -l mylink                      check the hard link
```

```
chen@chen-VirtualBox:~$ ln dir/data.txt mylink
chen@chen-VirtualBox:~$ cat mylink
1234567890
chen@chen-VirtualBox:~$ ls -l mylink
-rw-rw-r-- 2 chen chen 11 Feb 18 12:44 mylink
```

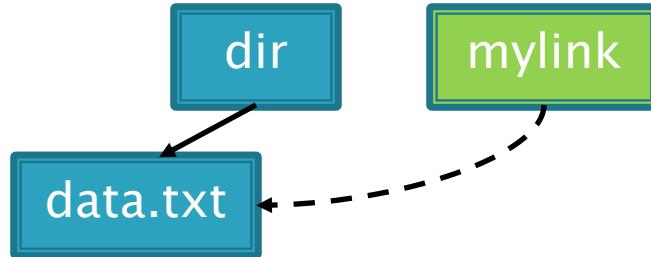
Example in Bash: Symbolic links

code/lecture_4_code/try

```
$ ln -s dir/data.txt mylink  
$ cat mylink  
1234567890
```

create a symbolic link

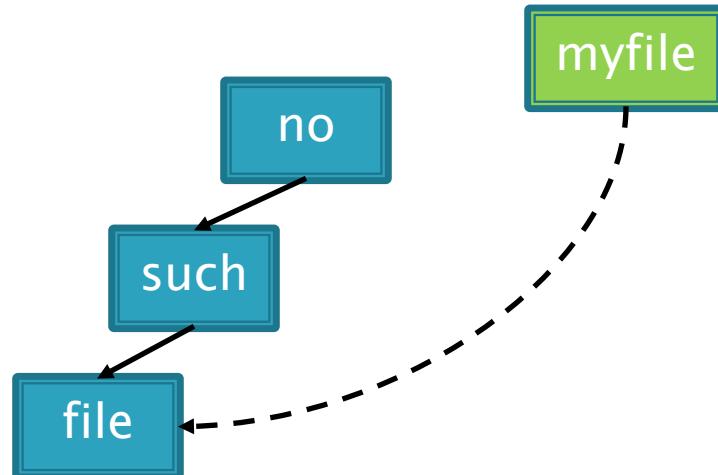
show the content in data.txt



Example in Bash: Symbolic links

code/lecture_4_code/try

```
$ ln -s /no/such/file myfile          create a symbolic link  
$ ls myfile  
myfile  
$ cat myfile  
cat: myfile: No such file or directory  
$ ls -l myfile                         ls says it's there  
lrwxrwxrwx 1 sar 13 Jan 22 00:26 myfile -> /no/such/file      so we try to look at it  
try -l option
```



Example: Symbolic links

code/lecture_4_code/symlink

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <errno.h>
#include <string.h>
//Create a link to the dir/data.txt and read it through readlink()
int main(void){
    char buf[1024];
    char content[1024];
    ssize_t len;
    if(symlink("dir/data.txt", "mylink") != 0)
        printf("Error: %s\n", strerror(errno));
    if ((len = readlink("mylink", buf, sizeof(buf)-1)) != -1)
        buf[len] = '\0';
    printf("The file path: %s\n", buf);
    FILE * pFile;
    pFile = fopen ("mylink", "r");
    fread(content, 10, 1, pFile);
    printf("%s\n", content); //print out file content
    fclose(pFile);
    return 0,
}
```

% ./main

The file path: dir/data.txt
1234567890

Treatment of symbolic links by various functions

Function	Does not follow symbolic link	Follows symbolic link
access		•
chdir		•
chmod		•
chown		•
creat		•
exec		•
lchown	•	
link		•
lstat	•	
open		•
opendir		•
pathconf		•
readlink	•	
remove	•	
rename	•	
stat		•
truncate		•
unlink	•	

Example: Hard links

code/lecture_4_code/hardlink

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <errno.h>
#include <string.h>
                           //Create a hard link and then unlink the file
int main(void){
    char c;
    FILE *fp;
    if (link("dir/data.txt", "myhardlink") == 0) //create a hard link
        printf("Hard link succeeds\n");
```

Example: Hard links

code/lecture_4_code/hardlink

```
fp = fopen("myhardlink","r");
printf("File content:");
c = fgetc(fp);
while(c!=EOF)
{
    //print the content of the file
    printf ("%c",c);
    c = fgetc(fp);
}
fclose(fp);
printf("\n");
sleep(3);
if (unlink("myhardlink") < 0)
    printf("unlink error");
printf("file unlinked\n");
return 0;
```

% ./main

```
Hard link succeeds
File content:1234567890
file unlinked
```

//unlink the file

```
In dir/data.txt mylink
cat mylink
ls -l mylink
```

Homework Readings

- ▶ **Reading:**
 - CSAPP: Section 10.5–10.10
 - APUE: Section 3.10, Section 4