

# 456: Network-Centric Programming

Yingying (Jennifer) Chen

Web: <https://www.winlab.rutgers.edu/~yychen/>  
Email: [yingche@scarletmail.rutgers.edu](mailto:yingche@scarletmail.rutgers.edu)


Department of Electrical and Computer Engineering  
Rutgers University

# TCP UDP SCTP

TCP: Transmission Control Protocol

UDP: User Datagram Protocol

SCTP: Stream Control Transmission Protocol



# TCP (Transmission Control Protocol )

- ▶ TCP is a reliable, ordered, and error-checked protocol.
- ▶ TCP provides a stream of bytes between applications running on hosts communicating via an Internet Protocol (IP) network.
- ▶ TCP is a connection-oriented protocol, which means a connection is established and maintained until the application programs at each end have finished exchanging messages.

# UDP (User Datagram Protocol)

- ▶ UDP is used primarily for establishing low-latency and loss-tolerating connections between applications on the internet.
- ▶ Send short messages called datagrams
- ▶ It is an unreliable, connectionless protocol.

# SCTP (Stream Control Transmission Protocol)

- ▶ SCTP is a protocol for transmitting multiple streams of data at the same time between two end points that have established a connection in a network. Sometimes referred to as "next generation TCP".
- ▶ SCTP's multi-streaming allows data to be delivered in multiple, independent streams, so that if there is data loss in one stream, delivery will not be affected for the other streams.

# Sockets API

- General interface for network programming and inter-process communication
- A socket is a communication endpoint, a tap into the network
- Network protocol independent but usually used with Internet protocols
  - Allows variety of addressing formats
- Datagram
  - Unordered message-oriented communication
  - Application multiplexing
  - Usually mapped to UDP
- Stream
  - Application multiplexing
  - Reliable, flow controlled data stream
  - Usually mapped to TCP

# Creating a socket

- ▶ `int socket(int domain, int type, int protocol)`
  - Returns socket file descriptor

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

Returns: nonnegative descriptor if OK, -1 on error

- ▶ **domain:** specifies the protocol family
- ▶ **type:** socket type
- ▶ **protocol:** set to a specific protocol type or 0  
(select default AF\_UNIX)

# Creating a socket

- ▶ `int socket(int domain, int type, int protocol)`
- ▶ `domain`: selects the protocol family which will be used for communication.

Domain	Description
<code>AF_INET</code>	IPv4 protocols
<code>AF_INET6</code>	IPv6 protocols
<code>AF_LOCAL</code>	Unix domain protocols
<code>AF_ROUTE</code>	Routing sockets
<code>AF_KEY</code>	Key socket

# Creating a socket

- ▶ `int socket(int domain, int type, int protocol)`
- ▶ `type`: socket type, it specifies the communication semantics.

type	Description
SOCK_STREAM	stream socket
SOCK_DGRAM	datagram socket
SOCK_SEQPACKET	sequenced packet socket
SOCK_RAW	raw socket

**SOCK\_STREAM** Provides sequenced, reliable, two-way, connection-based byte streams.

**SOCK\_DGRAM** Supports datagrams (connectionless, unreliable messages of a fixed maximum length).

**SOCK\_RAW** Provides raw network protocol access (you can use your own structures for UDP and TCP packet headers)

**SOCK\_SEQPACKET** gives you the guarantees of SOCK\_STREAM (i.e., preservation of ordering, guaranteed delivery, no duplication), but with delineated packet boundaries just like SOCK\_DGRAM. So, basically it's a mix of the two protocol types.

In the TCP/IP-family, SCTP implements both SOCK\_STREAM (TCP-like) and SOCK\_SEQPACKET.

# Creating a socket

- ▶ `int socket(int domain, int type, int protocol)`
- ▶ `protocol`: Setting protocol to 0 to select system's default based on the given domain and type.


Protocol	Description
IPPROTO_TCP	TCP transport protocol
IPPROTO_UDP	UDP transport protocol
IPPROTO_SCTP	SCTP transport protocol

TCP: Transmission Control Protocol


UDP: User Datagram Protocol

SCTP: Stream Control Transmission Protocol

# Typical Socket Implementation of Client–Server Scenario



# Socket functions for elementary TCP client/server



# Network prefix and Host number

- ▶ IP address (v4) is 4 bytes (32bits) long.
- ▶ The network prefix identifies a network and the host number identifies a specific host (actually, interface on the network).



- ▶ How do we know how long the network prefix is?
  - The network prefix used to be implicitly defined (**class-based addressing, A,B,C,D...**)


Class	Left-most Bit	Starting IP Address	Last IP Address
A	0xxx	0.0.0.0	127.255.255.255
B	10xx	128.0.0.0	191.255.255.255
C	110x	192.0.0.0	223.255.255.255
D	1110	224.0.0.0	239.255.255.255
E	1111	240.0.0.0	255.255.255.255

- The network prefix now is flexible and is indicated by a **prefix/netmask (classless)**.


# Example

**Example:** www.example.com

- ▶ IP address is 128.143.137.144
- ▶ Using Prefix notation IP address is: **128.143.137.144/16**
  - Network prefix is **16** bits long
- ▶ Network mask is: **255.255.0.0**
  - > **Network id** (IP address **AND** Netmask) is: **128.143.0.0**
  - > **Host number** (IP address **AND** inverse of Netmask **0.0.255.255**) is: **137.144**



# Which IP address?



Internet Assigned Numbers Authority (IANA) reserves these for private IP addresses

Private IP address space	
From	To
10.0.0.0	10.255.255.255
172.16.0.0	172.31.255.255
192.168.0.0	192.168.255.255

# Internet Assigned Numbers Authority (IANA) Maintains Port Numbers (16bits)

IANA well-known / Privileged	IANA registered	IANA dynamic or private
1 1023	1024	49151 49152 65535
<ul style="list-style-type: none"><li>■ Common applications</li><li>■ E.g.<ul style="list-style-type: none"><li>□ Daytime 13</li><li>□ http 80</li></ul></li><li>■ Root access only</li></ul>	<ul style="list-style-type: none"><li>■ Registered, less common applications</li><li>■ E.g.<ul style="list-style-type: none"><li>□ X Window server</li></ul></li></ul>	<ul style="list-style-type: none"><li>■ Used for private applications</li><li>■ Used as ephemeral ports</li></ul>

TCP, UDP, ... have separate sets of port numbers

# Useful commands

## ▶ Network statistics

- netstat -i

Displays network interfaces and their statistics

```
chen@chen-VirtualBox:~$ netstat -i
```

Kernel Interface table											Flg
Iface	MTU	Met	RX-OK	RX-ERR	RX-DRP	RX-OVR	TX-OK	TX-ERR	TX-DRP	TX-OVR	Flg
enp0s3	1500	0	44849	0	0	0	32044	0	0	0	B
MRU											
lo	65536	0	10873	0	0	0	10873	0	0	0	L
RU											

- ▶ Iface (Name of the interface)
  - Primary network card “eth0”
  - The loopback interface “lo” (a virtual network card that allows the computer to make networking connections to itself)
- ▶ MTU (Maximum Transmission Unit this interface can send at a time (in byte))
- ▶ Flg
  - B: broadcast capability; M: Multicast capability; L loopback interface

enp0s3 "ethernet network peripheral # serial #"  
Ubuntu virtual machine

# Useful commands

```
chen@chen-VirtualBox:~$ netstat -i
Kernel Interface table
Iface      MTU Met     RX-OK RX-ERR RX-DRP RX-OVR     TX-OK TX-ERR TX-DRP TX-OVR Flg
enp0s3      1500 0      44849      0      0 0          32044      0      0      0 B
MRU
lo         65536 0      10873      0      0 0          10873      0      0      0 L
RU
```

## ▶ RX-OK/ERR/DRP/OVR

- Statistics about the packets that have been received by the interface
- OK: correctly received
- ERR: received but with incorrect checksum
- DRP: dropped because the receiver buffer is full
- OVR: dropped because the kernel couldn't get to it in time


TX-OK/ERR/DRP/OVR show similar statistics at the transmitter

# Useful commands

- ▶ ifconfig <interface>
  - Show IP address and configuration of interface
- ▶ Note the */o* loopback interface -> good choice for your assignments
  - The loopback device is a special, virtual network interface that your computer uses to communicate with itself.
  - 127.0.0.1; localhost

```
-VirtualBox:~$ ifconfig lo
lo      Link encap:Local Loopback
        inet addr:127.0.0.1  Mask:255.0.0.0
        inet6 addr: ::1/128 Scope:Host
              UP LOOPBACK RUNNING  MTU:65536  Metric:1
              RX packets:70 errors:0 dropped:0 overruns:0 frame:0
              TX packets:70 errors:0 dropped:0 overruns:0 carrier:0
              collisions:0 txqueuelen:1000
              RX bytes:6402 (6.4 KB)  TX bytes:6402 (6.4 KB)
```

# TCP Three-Way Handshake (Connection Establishment)



# TCP Three-Way Handshake (Connection Establishment)

- ▶ The client issues an active open by calling *connect()*. This causes the client TCP to send a "**synchronize**" (**SYN**) segment, which tells the server the client's initial sequence number for the data that the client will send on the connection.
  - Normally, there is no with the SYN; it just contains an IP header, a TCP header, and possible TCP options.
- ▶ The server must **acknowledge (ACK)** the client's SYN and the server must also send its own SYN containing the initial sequence number for the data that the server will send on the connection.
  - The server sends its SYN and the ACK of the client's SYN in a single segment.

# How do client and server know which port to choose?

# bind()

```
#include <sys/socket.h>

int bind(int sockfd, struct sockaddr *my_addr, int addrlen);
```

Returns: 0 if OK, -1 on error

- Called by a TCP server
- The bind() function tells the kernel to associate the server's socket address in *my\_addr* with the socket descriptor *sockfd*.  
The *addrlen* argument is sizeof(sockaddr\_in)
- Assigns a local address to a socket (e.g. IP address and/or port number)

# listen()

- Called by a TCP server
- After socket() and bind() but must be before calling accept()

```
#include <sys/socket.h>

int listen(int sockfd, int backlog);
```

Returns: 0 if OK, -1 on error


- Backlog specifies the buffer size (queue length) for incoming connections
  - We will typically set it to a large value, such as 1024.
- Configures a socket to start accepting incoming connections
- The kernel maintains two queues:
  - Incomplete connection queue
  - completed connection queue

# The two queues maintained by TCP for a listening socket

- ▶ An incomplete connection queue
  - Contains an entry for each SYN packet that has arrived from a client for which the server is awaiting completion of the TCP three-way handshake.
  - The socket is in the SYN\_RCVD state
- ▶ A completed connection queue
  - Contains an entry for each client with whom the TPC three-way handshake has completed
  - The socket is in the ESTABLISHED state

# The two queues maintained by TCP for a listening socket

Packets exchange during the connection establishment with these two queues



# Connecting to a server

- ▶ `connect()`: a TPC client to establish a connection with a TCP server

```
#include <sys/socket.h>

int connect(int sockfd, struct sockaddr *serv_addr, int addrlen);
```

Returns: 0 if OK, -1 on error

- ▶ Requires open socket and destination address
- ▶ Several possible errors:
  - **ETIMEDOUT** – no response received after connection attempt
  - **ECONNREFUSED** – the server has refused the connection attempt (often wrong IP or port number)
  - **EHOSTUNREACH** – a router has notified the client that the destination could not be found

# accept()

- Called by a TCP server to wait for the request from a client


```
#include <sys/socket.h>

int accept(int listenfd, struct sockaddr *addr, int *addrlen);
```


Returns: nonnegative connected descriptor if OK, -1 on error

- The *listenfd* is the descriptor obtained from *socket()*  
Returns the next completed connection from the established queue
  - Return value is a different value for the connected socket: *connfd*
- Blocking
- addr* returns the address of the connected client
- addrlen* is a value-result argument (takes length of input structure, returns number of bytes stored in this structure)


# The roles of the listening and connected descriptors



1. *Server blocks in accept, waiting for connection request on listening descriptor listenfd.*




2. *Client makes connection request by calling and blocking in connect.*



3. *Server returns connfd from accept. Client returns from connect. Connection is now established between clientfd and connfd.*

# Handling multiple clients

server



socket()

bind(); listen();

accept() - *blocks*

Put in incomplete queue

Put in incomplete queue

accept() – *returns connfd x1 uses port x2*

Put in established queue

accept() – *returns connfd y1 uses port y2*

Put in established queue

# Socket Address Structures

---

*sockaddr: socketbits.h (included by socket.h), sockaddr\_in: netinet/in.h*

```
/* Generic socket address structure (for connect, bind, and accept) */
struct sockaddr {
    unsigned short   sa_family;      /* Protocol family */
    char             sa_data[14];    /* Address data.  */
};

/* Internet-style socket address structure */
struct sockaddr_in {
    unsigned short   sin_family;    /* Address family (always AF_INET) */
    unsigned short   sin_port;     /* Port number in network byte order */
    struct in_addr   sin_addr;     /* IP address in network byte order */
    unsigned char    sin_zero[8];   /* Pad to sizeof(struct sockaddr) */
};
```

---

*sockaddr: socketbits.h (included by socket.h), sockaddr\_in: netinet/in.h*

# Socket Address Structures

- ▶ The *sockaddr* is a *generic* socket address.
  - To accept any kind of socket address structure, we define sockets functions to expect a pointer to a generic *sockaddr* structure, and then require applications to cast pointers to protocol-specific structures to this generic structure.

```
struct sockaddr_in servaddr;  
....  
bind(listenfd, (sockaddr *) &servaddr, sizeof(servaddr));
```

- ▶ Internet socket addresses are stored in **16-byte** structures of the type *sockaddr\_in*. For Internet applications, the *sin\_family* member is *AF\_INET*, the *sin\_port* member is a 16-bit port number, and the *sin\_addr* member is a 32-bit IP address.

# Iterative Server (daytime server): lecture\_5\_code\daytime\daytimetcpsrv.c

```
#include <sys/socket.h>
#include <sys/types.h>
#include <time.h>
#include <strings.h>
#include <netinet/in.h>
#include <stdio.h>
#include <unistd.h>

#define MAXLINE          4096      /* max text line length */#define
LISTENQ            1024      /* 2nd argument to listen() */

typedef struct sockaddr SA;

int main(int argc, char **argv){

    int             listenfd, connfd;
    struct    sockaddr_in servaddr;
    char           buff[MAXLINE];
    time_t          ticks;
```

# Iterative Server (daytime server): lecture\_5\_code\daytime\daytimetcpsrv.c

```
listenfd = socket(AF_INET, SOCK_STREAM, 0);

bzero(&servaddr, sizeof(servaddr)); //zero a byte string
servaddr.sin_family = AF_INET;
servaddr.sin_addr.s_addr = htonl(INADDR_ANY);

/* allows the server to accept a client connection on any interface */

servaddr.sin_port = htons(13); /* daytime server, can only be run as root */

bind(listenfd, (SA *) &servaddr, sizeof(servaddr));

listen(listenfd, LISTENQ);

for ( ; ; ) {
    len = sizeof(cliaddr);
    connfd = accept(listenfd, (SA *) &cliaddr, &len);

    ticks = time(NULL);
    snprintf(buff, sizeof(buff), "%.24s\r\n", ctime(&ticks)); //Write formatted output to sized buffer
    write(connfd, buff, strlen(buff));

    close(connfd);
}
```

Terminal 1  
**\$ sudo ./server**

Terminal 2  
**\$ ./client 127.0.0.1**

The **htonl** function converts a 32-bit integer from host byte order to network byte order.  
The **htons** function performs corresponding conversions for 16-bit integers.

# Address Conversion

- ▶ Convert string representation (“192.168.0.2”) into binary representation
- ▶ “Presentation” to/from “Numeric”
  - `int inet_pton(int family, const char *strptr, void *addrptr);`
  - `const char *inet_ntop(int family, const void *addrptr, char *strptr, size_t len);`

```
#include <arpa/inet.h>
```

```
int inet_pton(int family, const char *strptr, void *addrptr);
```

Returns: 1 if OK, 0 if input not a valid presentation format, -1 on error

```
const char *inet_ntop(int family, const void *addrptr, char *strptr, size_t len);
```

Returns: pointer to result if OK, `NULL` on error

The family could `AF_INET`.

# Iterative Server (daytime client): lecture\_5\_code\daytime\daytimetcpccli.c

```
#include <sys/socket.h>
#include <sys/types.h>
#include <time.h>
#include <strings.h>
#include <netinet/in.h>
#include <stdio.h>
#include <unistd.h>
#include <errno.h>
#include <dirent.h>
#include <stdlib.h>
#include <stdarg.h>
#include <arpa/inet.h>

#define MAXLINE          4096      /* max text line length */#define

typedef struct sockaddr SA;
```

# Iterative Server (daytime client): lecture\_5\_code\daytime\daytimetcpccli.c

```
int main(int argc, char **argv) {
    int sockfd, n;
    char recvline[MAXLINE + 1];
    struct sockaddr_in servaddr;

    if (argc != 2)
        err_quit("usage: a.out <IPaddress>");

    if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        err_sys("socket error");

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    servaddr.sin_port = htons(13); /* daytime server */
    if (inet_pton(AF_INET, argv[1], &servaddr.sin_addr) <= 0)
        err_quit("inet_pton error for %s", argv[1]);
    if (connect(sockfd, (SA *) &servaddr, sizeof(servaddr)) < 0)
        err_sys("connect error")
```

# Iterative Server (daytime client): lecture\_5\_code\daytime\daytimetcpccli.c

```
while ( (n = read(sockfd, recvline, MAXLINE)) > 0) {  
    recvline[n] = 0; /* null terminate */  
    if (fputs(recvline, stdout) == EOF)  
        err_sys("fputs error");  
}  
if (n < 0)  
    err_sys("read error");  
}
```

# Server (hellomsg server): lecture\_5\_code\hellomsg\server.c

```
#include <stdio.h>
#include <sys/socket.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <string.h>
#include <unistd.h>
#define PORT 8080
{
    int server_fd, new_socket, valread;
    struct sockaddr_in address;
    int opt = 1;
    int addrlen = sizeof(address);
    char buffer[1024] = {0};
    char *hello = "Hello from server";
    //Creating socket file descriptor
    if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0)
    {
        perror("socket failed");
        exit(EXIT_FAILURE);
    }
```

In this code, the client sends a hello message to the server and the server returns a hello message.

Terminal 1

\$ ./server

Terminal 2

\$ ./client

# Server (hellomsg server): lecture\_5\_code\hellomsg\server.c

```
address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons( PORT );

        // bind socket to the port 8080
if (bind(server_fd, (struct sockaddr *)&address, sizeof(address))<0)
{
    perror("bind failed");
    exit(EXIT_FAILURE);
}
if (listen(server_fd, 1024) < 0)
{
    perror("listen");
    exit(EXIT_FAILURE);
}
if ((new_socket = accept(server_fd, (struct sockaddr *) &address, (socklen_t*)&addrlen))<0)
{
    perror("accept");
    exit(EXIT_FAILURE);
}
```

# Server (hellomsg server): lecture\_5\_code\hellomsg\server.c

```
valread = read( new_socket , buffer, 1024);
printf("%s\n",buffer );
write(new_socket , hello , strlen(hello));
printf("Hello message sent\n");
return 0;
}
```

# Client (hellomsg client): lecture\_5\_code\hellomsg\client.c

```
// Client side
#include <stdio.h>
#include <sys/socket.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#define PORT 8080

int main(int argc, char const *argv[])
{
    struct sockaddr_in address;
    int sock = 0, valread;
    struct sockaddr_in serv_addr;
    char *hello = "Hello from client";
    char buffer[1024] = {0};
```

# Client (hellomsg client): lecture\_5\_code\hellomsg\client.c

```
if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0)
{
    printf("\n Socket creation error \n");
    return -1;

memset(&serv_addr, '0', sizeof(serv_addr));
serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(PORT);

// Convert IP addresses from text to binary form
if(inet_pton(AF_INET, "127.0.0.1", &serv_addr.sin_addr)<=0)
{
    printf("\nInvalid address/ Address not supported \n");
    return -1;
}
```

# Client (hellomsg client): lecture\_5\_code\hellomsg\client.c

```
if (connect(sock, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0)
{
    printf("\nConnection Failed \n");
    return -1;
}
write(sock , hello , strlen(hello));
printf("Hello message sent\n");
valread = read( sock , buffer, 1024);
printf("%s\n",buffer );
return 0;
}
```

# Homework Readings

- Readings:

- ❑ Unix Network Programming: Section 3,4
  - ❑ Computer Systems – A Programmer’s Perspective:  
Section 10.4 + 11.4