

456: Network-Centric Programming

Yingying (Jennifer) Chen

Web: <https://www.winlab.rutgers.edu/~yychen/>
Email: yingche@scarletmail.rutgers.edu

Department of Electrical and Computer Engineering
Rutgers University

Working with Unix Processes

Processes

- ▶ Recall: Process is a running instance of a program
 - Multiple instances of the same program can run at the same time
 - E.g., open two terminal windows to run the terminal program twice
- ▶ Operating systems use the concept of “multitask” to give the impression that all processes run at the same time

Processes

- ▶ Try the “ps” command to see a list of active processes

```
chen@chen-VirtualBox:~$ ps
  PID TTY          TIME CMD
15101 pts/17    00:00:00 bash
15627 pts/17    00:00:00 ps
```

- ▶ Two processes are shown
 - Bash: the shell running on this terminal
 - The running instance of the “ps” program

Note: ps is short for “process status”

Processes

- Try the “ps” command to see a list of active processes

```
chen@chen-VirtualBox:~$ ps
  PID TTY          TIME CMD
15101 pts/17    00:00:00 bash
15627 pts/17    00:00:00 ps
```

- Four items are displayed
 - PID: the identifier of the process
 - TTY (terminal type): the name of the console or terminal
 - pts (pseudo terminal slave)
 - TIME: the amount of CPU time the process has utilized
 - CMD: the name of the command that launched the process

Options of the ps Command

Option	Description
-a	Displays all processes on a terminal, with the exception of group leaders.
-c	Displays scheduler data.
-d	Displays all processes with the exception of session leaders.
-e	Displays all processes.
-f	Displays a full listing.
-g/ <i>list</i>	Displays data for the <i>list</i> of group leader IDs.
-j	Displays the process group ID and session ID.
-l	Displays a long listing
-o	Specify the output format (e.g., ps -o pid,user,start_time,cmd)
-p/ <i>list</i>	Displays data for the <i>list</i> of process IDs.
-s/ <i>list</i>	Displays data for the <i>list</i> of session leader IDs.
-t/ <i>list</i>	Displays data for the <i>list</i> of terminals.
-u/ <i>list</i>	Displays data for the <i>list</i> of usernames

```
chen@chen-VirtualBox:~$ ps -a
 PID TTY      TIME CMD
 2109 pts/1    00:00:00 ps
```

fork

- System call to create processes
 - Clone (and run) another instance of a process (parent process)
- Create a process using **pid_t fork(void)**
 - Create a child process (a duplicate process)
 - A child process executes the **same program** from the same place
 - Typedef **pid_t** to refer to process IDs (defined in `<sys/types.h>`)

fork

Why?

- Return twice (for parent and child)
 - For the parent process: the return value is the process ID of the child
 - For the child process, the return value is zero
- The return value
 - Is used to check whether the program is running as the parent or the child process

init process

How is the first process created?

- ▶ Every process has an identifier (PID)
- ▶ “init” is the very first process to run after booting (PID 1)
 - Created by the kernel
 - parent of all the processes
 - PID 0: kernel scheduler
- ▶ Every process except init has a parent
 - Parent is the process that calls fork
- ▶ init adopts processes that lose their parent (“orphan process”)

Example (Using fork to duplicate a process): lecture_7_code/fork_example/fork_example.c

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    printf ("the main program process ID is %d\n", (int) getpid ());
    pid = fork();
    if (pid != 0) {
        printf ("this is the parent process, with id %d\n", (int) getpid ());
        printf ("the child's process ID is %d\n", (int) pid);
    }
    else
        printf ("this is the child process, with id %d\n", (int) getpid ());
    return 0;
}
```

Fork system
Call –
“Returns twice”

Terminal
\$./main

Example (**Concurrent server**): lecture_7_code/concurrent_server/server.c

```
#include <stdio.h>
#include <sys/socket.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <string.h>
#include <unistd.h>
#define PORT 8080
int main(int argc, char const *argv[])
{
    int server_fd, new_socket, valread;
    struct sockaddr_in address;
    int addrlen = sizeof(address);
    char buffer[1024] = {0};
    int pid;
    char *response = "Response from server";
```

Example (**Concurrent server**): lecture_7_code/concurrent_server/server.c

```
if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0){  
    perror("socket failed");  
    exit(EXIT_FAILURE);  
}  
address.sin_family = AF_INET;  
address.sin_addr.s_addr = htonl (INADDR_ANY);  
address.sin_port = htons( PORT );      //attaching socket to the port 8080  
if (bind(server_fd, (struct sockaddr *)&address,  
          sizeof(address))<0) {  
    perror("bind failed");  
    exit(EXIT_FAILURE);  
}  
if (listen(server_fd, 1024) < 0) {  
    perror("listen");  
    exit(EXIT_FAILURE);  
}
```

Example (Concurrent server): lecture_7_code/concurrent_server/server.c

```
while(1){  
    if ((new_socket = accept(server_fd, (struct sockaddr *)&address, (socklen_t*)&addrlen))<0) {  
        perror("accept");  
        exit(EXIT_FAILURE);  
    }  
  
    /*Create a new process to cope with each new client request*/  
    if ( (pid = fork()) == 0) {  
        /* child process */  
        close(server_fd);  
        /* close listening socket */  
        /* Cope with the request */  
        valread = read( new_socket , buffer, 1024); // receive the msg  
        printf("%s\n",buffer );  
        write(new_socket, response, strlen(response)); //send the response from server  
        printf("Response msg sent\n");  
        exit(0);  
    }  
  
    close(new_socket); /* parent closes connected socket */  
    printf("after close socket\n");  
}  
return 0;  
}
```

Example (**Concurrent server**): lecture_7_code/concurrent_server/client.c

```
#include <stdio.h>
#include <sys/socket.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#define PORT 8080

int main(int argc, char const *argv[])
{
    struct sockaddr_in address;
    int sock1, sock2, valread;
    struct sockaddr_in serv_addr;
    char *hello1 = "Hello from client 1";
    char *hello2 = "Hello from client 2";
    char buffer[1024] = {0};
```

Example (**Concurrent server**): lecture_7_code/concurrent_server/client.c

```
if ((sock1 = socket(AF_INET, SOCK_STREAM, 0)) < 0)
{
    printf("\n Socket creation error \n");
    return -1;
}
if ((sock2 = socket(AF_INET, SOCK_STREAM, 0)) < 0)
{
    printf("\n Socket creation error \n");
    return -1;
}
// initialize serv_addr
memset(&serv_addr, '0', sizeof(serv_addr));

serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(PORT);
```

Example (**Concurrent server**): lecture_7_code/concurrent_server/client.c

```
if/inet_pton(AF_INET, "127.0.0.1", &serv_addr.sin_addr)<=0) {  
    printf("\nInvalid address/ Address not supported \n");  
    return -1;  
}  
if (connect(sock1, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {  
    printf("\nConnection Failed \n");  
    return -1;  
}  
// send messages via two sockets: sock1  
write(sock1 , hello1 , strlen(hello1));           //send msg via sock1  
printf("Hello message sent from client 1\n");  
valread = read(sock1, buffer1, 1024);             //receive msg from server  
printf("%s\n", buffer1);
```

Example (**Concurrent server**): lecture_7_code/concurrent_server/client.c

```
if (connect(sock2, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0){  
    printf("\nConnection Failed \n");  
    return -1;  
}  
// send messages via two sockets: sock2  
write(sock2 , hello2 , strlen(hello2));           //send msg via sock2  
printf("Hello message sent from client 2\n");  
valread = read(sock2, buffer2, 1024);             //receive msg from server  
printf("%s\n", buffer2);  
return 0;  
}
```

Terminal 1
\$./server

Terminal 2
\$./client

Running a Different Program in the New Process

- **fork**: Make a child process an exact copy of its parent process
- **exec**: Execute a new program image in a new process
- Two steps:
 - Use fork() to create a new process
 - Use **exec family** function to load the new program into that process
 - The base exec followed by one or more letters

Examples of functions in exec family:
execv, execve, execvp, execl, execle and execlp,

Running a Different Program in the New Process – exec family of function

- Loads new program into process
- Exec p/v/e
 - p: search current execution path for program named in the file argument
 - v: accept an argument list being passed to the new program as an array (vector) of pointers
 - e: an array of pointers to environment variables is explicitly passed to the new process image

Example:

```
execvp (new_program, arg_list);
```

Example (Using exec to run a different program): lecture_7_code/fork_and_exec/fork_and_exec.c

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
int spawn (char* program, char** arg_list){
pid_t child_pid;
/* Duplicate this process. */
child_pid = fork ();
if (child_pid != 0)           /* This is the parent process. */
    return child_pid;
else {
    /* Now execute the program, searching for it in the current path. */
execvp (program, arg_list);

    /* The execvp function returns only if an error occurs. */
fprintf (stderr, "an error occurred in execvp\n");
abort (); /*abnormal program termination*/
}
}
```

/* Spawn a child process running a new program.
program is the name of the program to run;
arg_list is a NULL-terminated list of character
strings to be passed as the program's argument
list. Returns the process ID of the spawned
process. */

Example (Using exec to run a different program): lecture_7_code/fork_and_exec/fork_and_exec.c

```
int main ()
{
    /* The argument list to pass is the “ls -l /”. */
    char* arg_list[] = {
        “ls”,           /* argv[0], the name of the program. */
        “-l”,           /* long format */
        “/”,            /* root directory */
        NULL /* The argument list must end with a NULL. */
    };
    printf(“print in main process before calling spawn.”);
    /* Spawn a child process running the “ls” command. */
    spawn (“ls”, arg_list);           /*The new process runs “ls” command*/
    printf (“done with main program\n”);
    return 0;
}
```

Terminal
\$./main

Process Termination

- ▶ **return, exit()**
 - Perform cleanup operations (fclose on stdio streams)
 - Call exit handlers (register with atexit(...))
 - **atexit()** registers the function to run at process termination
 - If multiple handlers are registered, called in **reverse order of registration**
- int atexit (void (*func)(void));
Return 0 if successfully
- ▶ **Other termination methods**
 - **_exit()**
 - Return to kernel immediately
 - **abort()**
 - Abnormal program termination

Example (**atexit example**): lecture_7_code/atexit_example/atexit_example.c

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>

static void my_exit1(void), my_exit2(void);

int main(void)
{
    if (atexit(my_exit2) != 0)
        err_sys("can't register my_exit2");
    if (atexit(my_exit1) != 0)
        err_sys("can't register my_exit1");
    if (atexit(my_exit1) != 0)
        err_sys("can't register my_exit1");
    printf("main is done\n");
    return(0);
}
```

Example (**Atexit example**):

lecture_7_code/atexit_example/atexit_example.c

```
static void my_exit1(void)
{
    printf("first exit handler\n");
}
static void my_exit2(void)
{
    printf("second exit handler\n");
}
```

Terminal

\$./main

Output?

Output

- `%./atexit_example`

```
main is done
first exit handler
first exit handler
second exit handler
```

wait / waitpid

- System calls used in the parent process
 - Wait for a child process to finish executing
 - Retrieve information about its child's termination
 - **Clean up resources** used by the child process
 - Return child process ID
- *wait* pid_t wait(int *status);
 - Blocking
 - Wait for any child process
- *waitpid* pid_t waitpid(pid_t pid, int *status, int options);
 - Wait for a specific child
 - Can be nonblocking (when use with other functions)
- Retrieve termination status of child process
 - Macros available to distinguish between normal and abnormal termination
 - E.g., WIFEXITED(status)
 - Returns True if status was returned for a normally terminated child
 - Macros defined in the header file <sys/wait.h>

Example (**Wait example**): lecture_7_code/wait_example/wait_example.c

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
/* Use wait to the previous fork and exec example*/
int spawn(char* program, char** arg_list)
{
    pid_t child_pid;
    /* Duplicate this process. */
    child_pid = fork();
    if (child_pid != 0)
        /* This is the parent process. */
        return child_pid;
    else {
        /* Now execute PROGRAM, searching for it in the path. */
        execvp(program, arg_list);

        /* The execvp function returns only if an error occurs. */
        fprintf(stderr, "an error occurred in execvp\n");
        abort();
    }
}
```

Example (**Wait example**): lecture_7_code/wait_example/wait_example.c

```
int main()
{
    int child_status;
    /* The argument list to pass is the "ls -l /". */
    char* arg_list[] = {
        "ls", /* argv[0], the name of the program. */
        "-l",
        "/",
        NULL /* The argument list must end with a NULL. */
    };
    printf("print in main process before calling spawn.");
    spawn("ls", arg_list); /* Spawn a child process running the "ls" command.

    wait(&child_status); /*wait for a child process to complete.*/
    if(WIFEXITED(child_status)) /* WIFEXITED returns a nonzero value if child terminates normally*/
        printf("the child process exited normally, with exit code %d\n",WEXITSTATUS(child_status));
    else
        printf("the child process exited abnormally\n");

    printf ("done with main program\n");
    return 0;
}
```

Terminal
\$./main

Note: WEXITSTATUS macro returns
the exit code specified by the child

Exit Code

- ▶ Every command returns an exit status (sometimes referred to as a return status or exit code).
- ▶ A successful command returns a 0
- ▶ Anything other than a 0 status is undesirable

exit code

- 1 - Catchall for general errors
- 2 - Misuse of shell builtins (according to Bash documentation)
- 126 - Command invoked cannot execute
- 127 - “command not found”
- 128 - Invalid argument to exit
- 128+n - Fatal error signal “n”
- 130 - Script terminated by Control-C
- 255* - Exit status out of range

Zombies

- A zombie process is a process that has terminated but has not been cleaned up yet
- It is the responsibility of the parent process to clean up its zombie children
- When the parent process calls wait, the zombie child's termination status is extracted and the child process is deleted
- If not being cleaned, the zombie process stays around in the system

Example (Making a zombie process):

lecture_7_code/zombie_process/zombie_process.c

```
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
int main ()
{
    pid_t child_pid;
    /* Create a child process. */
    child_pid = fork ();
    if (child_pid > 0) {
        sleep (60);           /* This is the parent process. Sleep for a minute. */
    }
    else {
        /* This is the child process. Exit immediately. */
        exit (0);
    }
    return 0;
}
```

Terminal 1

\$./main

Terminal 2

\$ ps -a -o pid,tty,stat,cmd

While the program is still running, list the processes in another window:

Example (Making a zombie process)

- Output:

```
PID TT      STAT CMD
3755 pts/17  S+    ./main
3756 pts/17  Z+    [main] <defunct>
3757 pts/18  R+    ps -a -o pid,tty,stat,cmd
```

- The child process is marked as <defunct> and the status “Z” means “zombie”

Process status code examples

- R Running or runnable (on run queue)
- S Interruptible sleep (waiting for an event to complete)
- Z Defunct ("zombie") process, terminated but not cleaned up by its parent.
- + is in the foreground process group

How can we pass information to a new process?

- Command Line Arguments
 - The process that can pass command-line arguments to the new program

Command Line Arguments – Output?

Example (**Echo all command-line arguments to standard output**):

lecture_7_code/echoarg/echoarg.c

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int i;

    for (i = 0; i < argc; i++)      /* echo all command-line args */
        printf("argv[%d]: %s\n", i, argv[i]);
    exit(0);
}
```

Terminal

\$./main arg1 TEST foo

Output

- \$./main arg1 TEST foo

```
argv[0]:./main  
argv[1]:arg1  
argv[2]:TEST  
argv[3]:foo
```

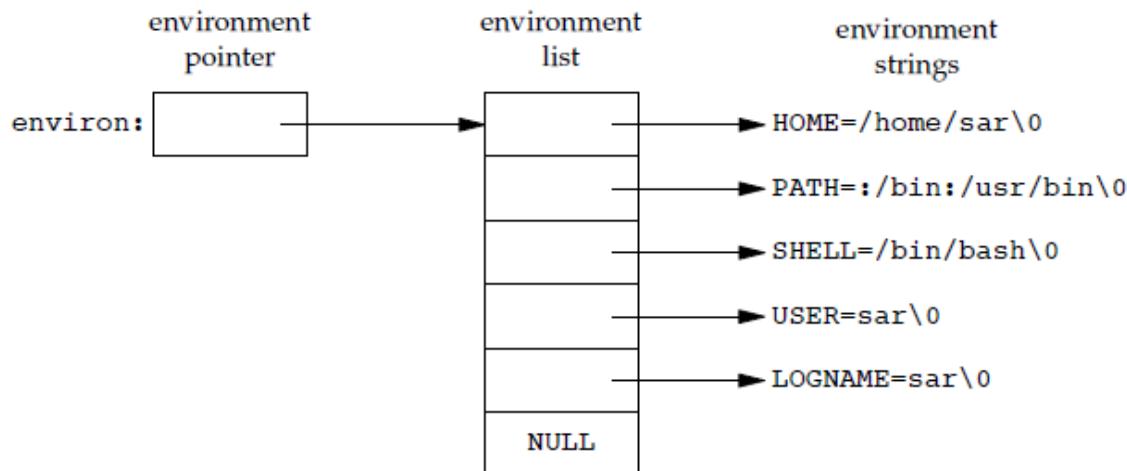
argv[0] stores the name of the program

How can we pass information to a new process?

- Environment list

Environment List

- ▶ An array of character pointers
- ▶ Each pointer contains the address of a null-terminated configuration string
- ▶ An example of the environment consisting of five strings
 - HOME contains the path to your home directory.
 - PATH contains a colon-separated list of directories through which Linux searches for commands you invoke.
 - SHELL The name of the user's default shell
 - LOGNAME the name of the user
 - USER contains your username, same purpose as LOGNAME



Environment List

- ▶ Accessible through global pointer variable
 - `extern char **environ;`
 - `char **` : NULL-terminated array of pointers to character strings
- ▶ Or more easily be accessed through:
 - `getenv` function in `<stdlib.h>`
`char* getenv(const char*)`

Example: `char* home = getenv ("HOME");`

Example (Print the execution environment):

lecture_7_code/printenv/printenv.c

```
#include <stdio.h>
/* The ENVIRON variable contains the environment. */
extern char** environ;
int main ()
{
    char** var;
    for (var = environ; *var != NULL; ++var)
        printf ("%s\n", *var);
    return 0;
}
```

Terminal
\$./main

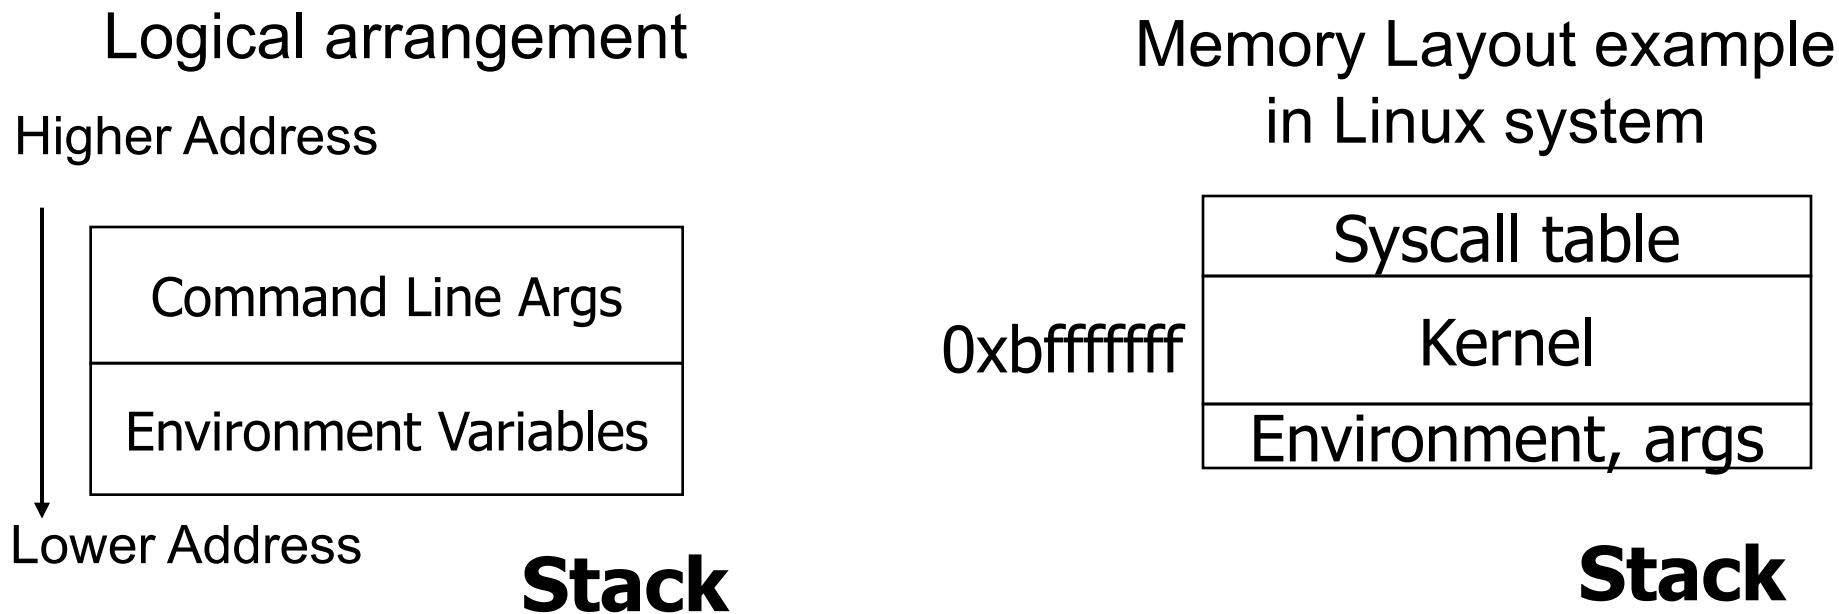
Memory Layout of a C Program

One C program consists of the following pieces:

- Text segment
 - Program Code
- Initialized data segment
 - Global variables
- Uninitialized data segment
 - (Initialized to 0 by kernel)
- Stack
 - Automatic variables (local variables)
 - The information saved when a function is called (e.g., return address)
- Heap
 - Dynamic memory allocation

Memory Layout of a C Program

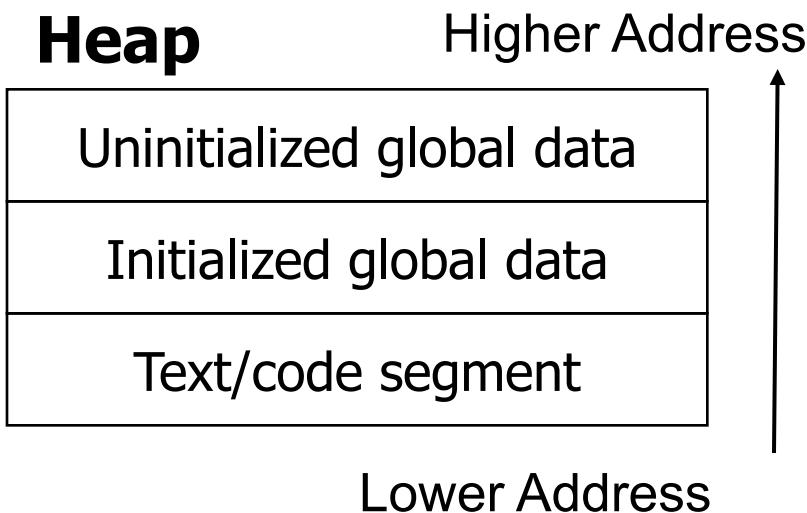
- **Stack segment:** memory area used by the process to store the local variables of function and other information that is saved every time a function is called, such as return address.



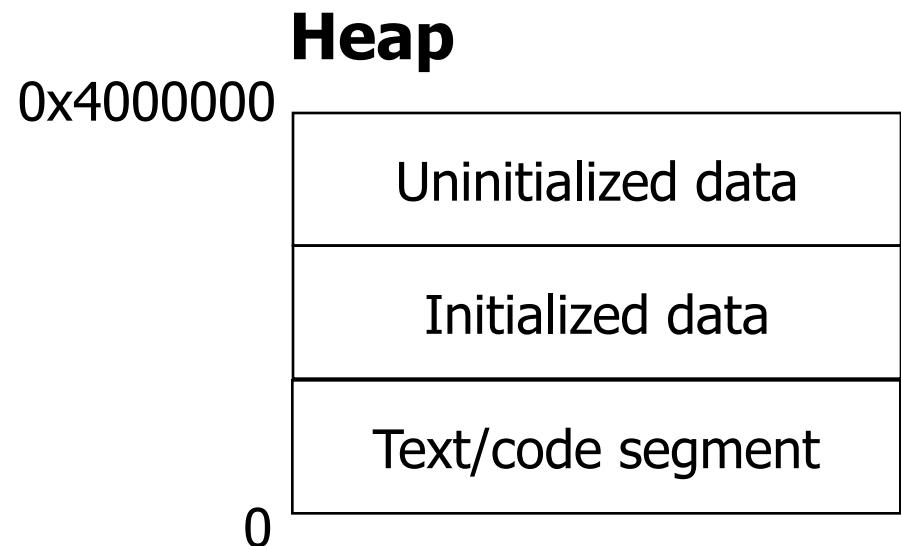
Memory Layout of a C Program

- **Heap segment** is used for dynamic memory allocation and shared among all the processes running in the system

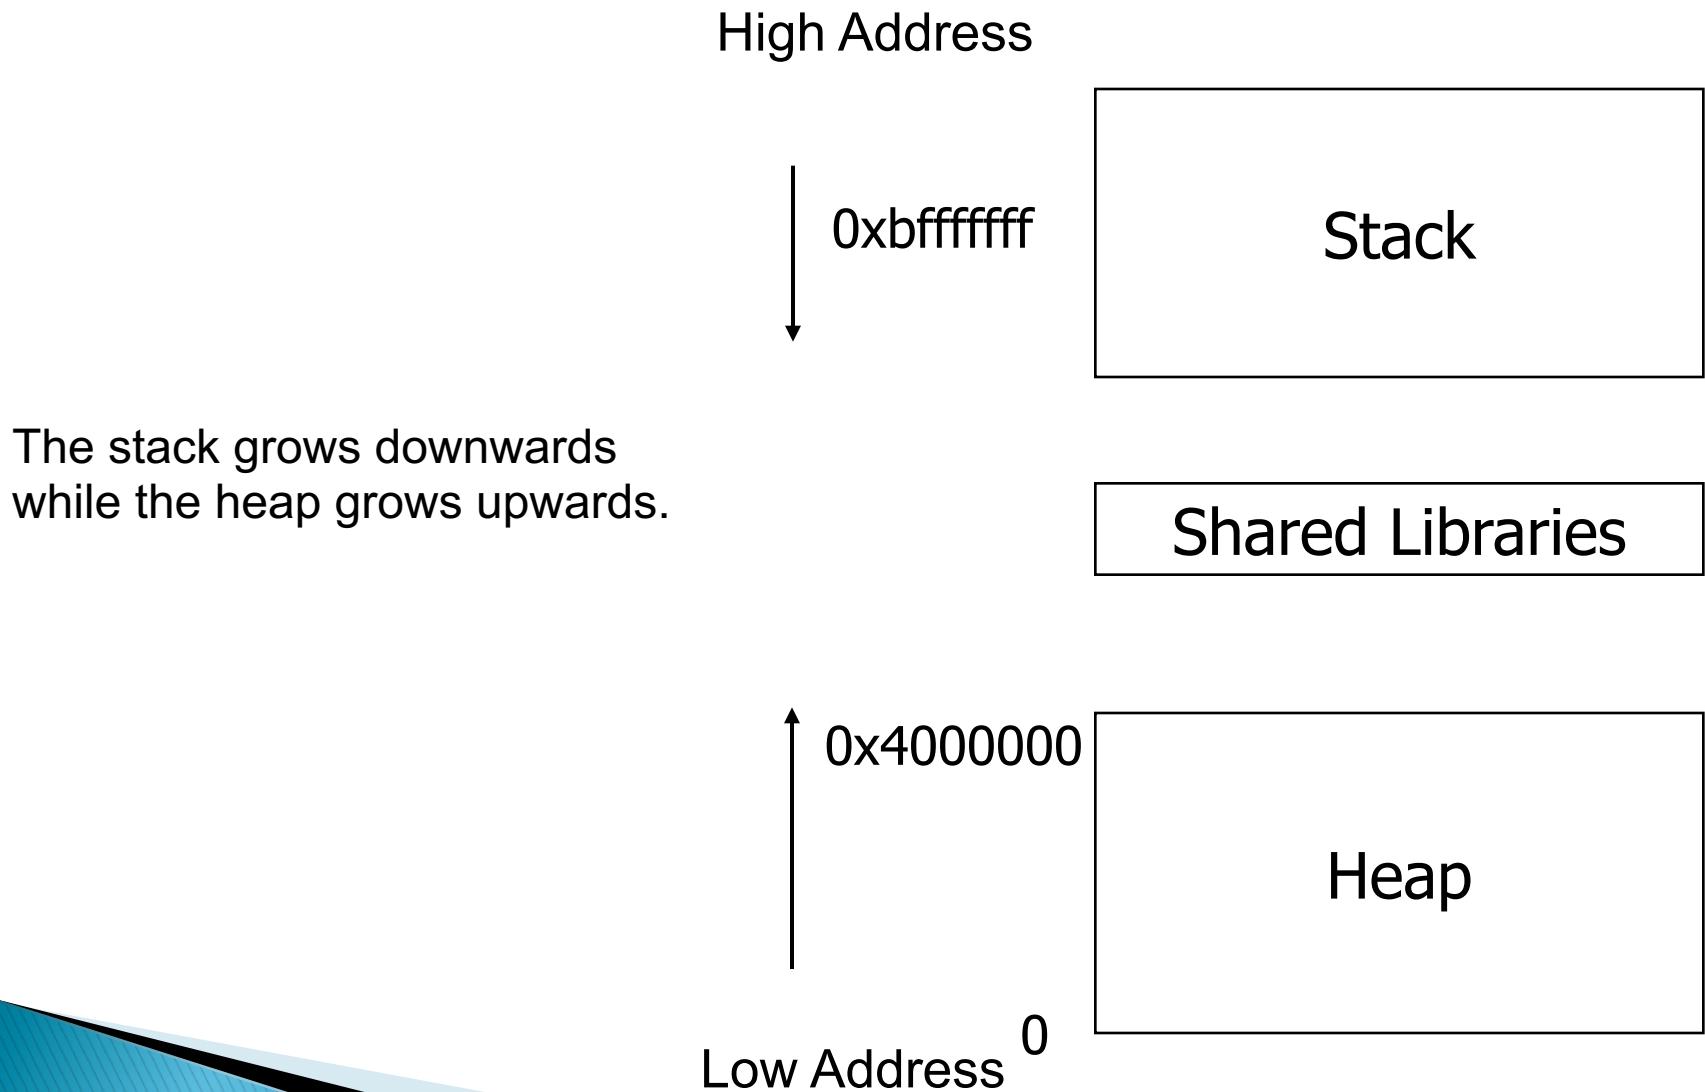
Logical arrangement



Memory Layout example
in Linux system



Memory Layout of a C Program



Example (Print variable order): lecture_7_code/printorder/main.c

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
int glob = 6;           /* external variable in initialized data */
char buf[] = "write to stdout\n";
void err_sys(const char* x);

/* This program demonstrates the fork function, showing how changes to variables in a
child process do not affect the value of the variables in the parent process*/

int main(void)
{
    int var;           /* automatic variable on the stack */
    pid_t pid;
    var = 88;
```

Example (Print variable order): lecture_7_code/printorder/main.c

```
if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1) //write function is not buffered
    err_sys("write error");
printf("before fork\n"); /* we don't flush stdout */

if ( (pid = fork()) < 0)
    err_sys("fork error");
else if (pid == 0) /* child */
    printf("the child process starts\n");
    glob++;
    /* modify variables */
    var++;
} else {
    sleep(2); /* parent */
    printf("this is the parent process\n");
}
printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
exit(0);
}

void err_sys(const char* x)
{
    perror(x);
    exit(1);
}
```

Output?

Terminal
\$./main

Output?

- \$./main

```
write to stdout
before fork
the child process starts
pid = 430, glob = 7, var = 89
this is the parent process
pid = 429, glob = 6, var = 88
```

- Last two lines may be reversed in some cases
(there is no rule whether parent or child will execute first)
- Is the output the same if we redirect to a file?

Output redirected to a file?

- \$./main > out.txt

```
write to stdout
before fork
the child process starts
pid = 430, glob = 7, var = 89
before fork
this is the parent process
pid = 429, glob = 6, var = 88
```

- Why?

Why?

- write function is not buffered (data written once)
- Stdout by default is line buffered when connected to terminal
 - End of line character flushes the buffer (e.g., “\n”)
- Stdout is fully buffered otherwise (e.g., directed to a file)
 - Child process inherits a copy of the buffer contents

Output redirected to a file

```
write to stdout
before fork
the child process starts
pid = 430, glob = 7, var = 89
before fork
this is the parent process
pid = 429, glob = 6, var = 88
```

Race Conditions

- When developing systems that involve multiple processes, we need to prevent race conditions
 - A race hazard (or race condition) is a flaw in a system or process where the output exhibits unexpected critical dependence on the relative timing of events.
 - For example, access shared memory at the same time or overwrite information before it is read.

Example (**Race condition**):

lecture_7_code/race_condition/race_condition.c

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
/*a function of printing a single character of the string at a time*/
static void charatatime(char *str);
void err_sys(const char* x);

Int main(void)
{
    pid_t pid;
    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) {
        charatatime("Child process wants to print a long string.\n");
    } else {
        charatatime(" Parent process is doing something else.\n");
    }
    exit(0);
}
```

Example (Race condition):

lecture_7_code/race_condition/race_condition.c

```
static void charatatime(char *str)
{
    char *ptr;
    int   c;

    setbuf(stdout, NULL);          /* set unbuffered */
    for (ptr = str; c = *ptr++; )
        putc(c, stdout);
}
```

```
void err_sys(const char* x)
{
    perror(x);
    exit(1);
}
```

Terminal
\$./main

Output:

Parent process is doingC hsiolmde tphrioncge sesl swea.n
ts to print a long string.

File Locking

File locks (“Record locks”)

- ▶ What happens when two process attempt to edit the same file at the same time ?
 - Each process independently maintains context information (such as write location), there may be coverage, crossover, and disordered content.
- ▶ **File locking** is the ability of a process to prevent other processes from modifying or reading a region of a file while the first process is working on it.

File locks (“Record locks”)

- ▶ Here are two types of locks :
 - Read lock (**shared** lock)
 - Any number of processes can have a shared read lock on a given region of a file.
 - All of them can read from the file but can't write to it.
 - Write lock (**exclusive** lock)
 - Only one process can have an exclusive lock on a given region of a file.
 - The other process are prevented from getting locks or reading and writing.

File locks (“Record locks”)

- Synchronize access from multiple processes to a shared file
- File locks provide fine-grained mechanisms:
 - Read and write locks: Multiple simultaneous reads are possible / Any write operation must be exclusive
 - Locking on byte ranges instead of entire file
- Use *fcntl()* with **locking commands** and **lock data structure** to acquire locks

Lock compatibility chart

	Read	Write
Read	Allowed	--
Write	--	--

fcntl()

- ▶ fcntl() function can change the properties of a file that is already open.
 - The fcntl() can be used for file lock purposes :

cmd	description
F_GETLK,F_SETLK, F_SETLKW	Get/set/lock file lock, it is discussed later.

```
#include <fcntl.h>
int fcntl (int filedes, int cmd, struct flock *flockptr);
```

fcntl()

- ▶ For record locking ,the third argument of `fcntl()` is a pointer to a **flock** structure :

```
struct flock {  
    short l_type;          /* F_RDLCK, F_WRLCK, F_UNLCK */  
    short l_whence;        /* SEEK_SET, SEEK_CUR, SEEK_END */  
    off_t l_start;         /* offset in bytes, relative to l_whence */  
    off_t l_len;           /* length, in bytes ; 0 means lock to EOF */  
    pid_t l_pid;           /* returned with F_GETLK */  
}
```

`l_pid` field is the process ID of the process holding the lock. It is filled in by calling `fcntl()` with the `F_GETLK` command, but is ignored when making a lock.

`fcntl()`

- ▶ Three commands of `fcntl` (record lock) :
 - `F_GETLK`
 - To check the lock described by *flockptr* :
 - The information pointed by *flockptr* is overwritten by an existing lock .
 - If no lock exists, *flockptr.l_type* = `F_UNLCK`
 - `F_SETLK`
 - Set the lock described by *flockptr*.
 - If lock is not possible, `fcntl()` will return `errno` = `EAGAIN` or `EACCES`.
 - `F_SETLKW`
 - A blocking version of `F_SETLK`. It will wait until the desired region released.

Example (**Lock and read a file**): lecture_8_code/lock_unlock/lock_unlock.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <errno.h>

/*The following example demonstrates how to place a lock on 10 bytes from the offset 100
in a file and remove it afterwards*/

int main(int argc, char *argv[])
{
    int fd;
    struct flock fl;
    char buf[20];
    printf ("Opening file.\n");                                W+
    fd = open("testfile.txt", O_RDWR); // open a file for reading and writing.
    if (fd == -1){
        perror("File open failed");
        exit(EXIT_FAILURE);
    }
```

Example (**Lock and read a file**): lecture_8_code/lock_unlock/lock_unlock.c

```
/* Make a non-blocking request to place a write lock on 10 bytes from the offset 100 in a
file */

fl.l_type = F_WRLCK;
fl.l_whence = SEEK_SET;
fl.l_start = 100;
fl.l_len = 10;

printf ("Locking the file.\n");
if (fcntl(fd, F_SETLK, &fl) == -1) { //set non-blocking block
    if (errno == EACCES || errno == EAGAIN) {
        printf("Already locked by another process\n");
        /* We can't get the lock at the moment */
    } else {
        perror("Handle unexpected error");
    }
} else {
    printf ("Lock was granted.\n");
```

Example (**Lock and read a file**): lecture_8_code/lock_unlock/lock_unlock.c

```
/* Perform I/O on bytes 100 to 109 of file */
lseek(fd, 100, SEEK_SET); /* seek to offset 100 from the start */
read(fd, buf, 10); /* read from 100 to 109 bytes */
buf[10] = '\0'; //terminates the 10-byte buf string
printf("The 100-109 bytes in the file are: %s\n", buf);
/*can also perform write process in the file*/

/* Unlock the locked bytes */
fl.l_type = F_UNLCK;
fl.l_whence = SEEK_SET;
fl.l_start = 100;
fl.l_len = 10;
printf ("Unlocking the file.\n");
if (fcntl(fd, F_SETLK, &fl) == -1){ //unlock the file
    perror("Handle unexpected error");
} else{
    printf ("Unlock was granted.\n");
}
exit(EXIT_SUCCESS)
```

Terminal:
\$./main

Example (**Wait lock and write**): lecture_8_code/lockfile/lockfile.c

```
#include <fcntl.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>

/*The program in Listing opens a file for writing, and then places a write lock
on it. The program waits for the user to hit Enter and then unlocks and closes
the file.*/
int main (int argc, char* argv[]){
    char* file = "testfile.txt";
    char* buffer = "The additional text we want to write to the file. ";
    int fd;
    struct flock lock;

    printf ("opening %s\n", file);
    /* Open a file descriptor to the file. */
    fd = open (file, O_WRONLY);
```

Example (**Wait lock and write**): lecture_8_code/lockfile/lockfile.c

```
/* Initialize the flock structure. */
memset (&lock, 0, sizeof(lock));
lock.l_type = F_WRLCK;

/* Place a write lock on the whole file. */
printf ("Place a blocking lock on the file\n");
fcntl (fd, F_SETLKW, &lock);

printf ("Write to the file\n");
lseek(fd, 0, SEEK_END); /* seek to the end of the file */
write(fd, buffer, strlen(buffer)); /* write the buffer to the end */
printf ("Finish writing; Hit Enter to unlock the file\n");

/* Wait for the user to hit Enter. */
getchar ();
printf ("Enter received, unlock the file now\n");
```

Example (**Wait lock and write**): lecture_8_code/lockfile/lockfile.c

```
/* Release the lock. */
lock.l_type = F_UNLCK;
fcntl (fd, F_SETLKW, &lock);
printf ("file is unlocked\n");
close (fd);
return 0;
}
```

1. Terminal 1:

\$./main
opening testfile.txt
Place a blocking lock on the file
Write to the file
Finish writing; Hit Enter to unlock the file

2. Terminal 2:

\$./main
opening testfile.txt
Place a blocking lock on the file

3. Terminal 1: hit Enter to unlock

4. Terminal 2: hit Enter to unlock

Homework Readings

- ▶ Advanced Linux Programming
 - Ch 2.1.1
 - Ch 2.1.6
 - Ch 3 (excluding 3.3)
- ▶ Advanced Programming in Unix Environment
 - Ch 7.1 – 7.6

Optional Homework

- ▶ One problem in each homework set
- ▶ Help to boost the points lost in the programming assignments
- ▶ Counted as 5 points in each homework