

456: Network-Centric Programming

Yingying (Jennifer) Chen

Web: <https://www.winlab.rutgers.edu/~yychen/>
Email: yingche@scarletmail.rutgers.edu

Department of Electrical and Computer Engineering
Rutgers University

Threads & Thread Synchronization

Threads vs. Processes

- ▶ Threads are “lightweight” processes
 - Represent an execution path through a program
- ▶ One process can contain many threads
 - If process terminates, all threads terminate
 - Process resources (e.g. file descriptors area shared among threads)
- ▶ Threads
 - Creation, context switching is faster
 - Share same address space

Advantages over Sequential Program

- ▶ Improve application responsiveness
 - Webserver answering requests
 - GUI responding to user input
- ▶ More efficient use of resources
 - Multiprocessor machines (e.g., Macbook Pro with 4-core CPU)
 - “Hyperthreading”
 - Intel's proprietary simultaneous multithreading (SMT) implementation used to improve parallelization of computations performed on x86 microprocessors.

Using Thread or Using Process

- ▶ A thread can do anything a process can do, and a thread could be considered a ‘lightweight’ process.
 - Since process can consist of multiple threads, a thread could be considered a ‘lightweight’ process.
- ▶ The essential difference between a thread and a process:
 - Threads are used for small tasks, whereas processes are used for more ‘heavyweight’ tasks – basically the execution of applications.

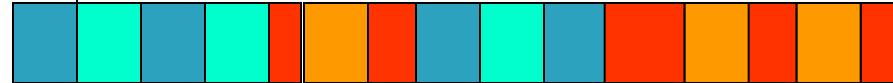
Interleaved Execution

Single-thread



Blocking I/O

Multi-thread

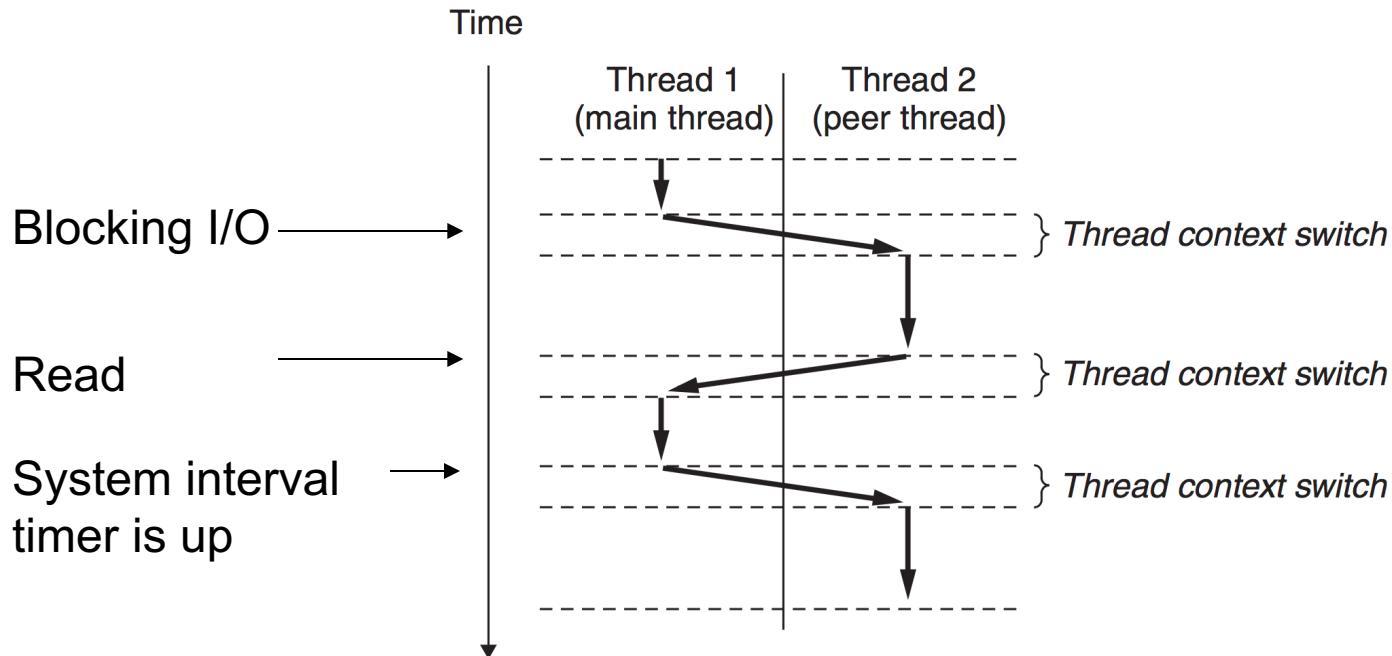


Execution Time
(Save the blocking time)

Thread Context Switch

- ▶ A procedure that CPU follows to change from one thread to another while ensuring no conflicts.
 - Store the state of a thread (e.g., registers and memory maps) so that it can be restored and execution resumed from the same point later
 - Pause the thread and resume another thread
 - Essential feature of a multitasking operating system
- ▶ When to Switch
 - Multi-threading
 - When a thread is interrupted

Concurrent Thread Execution



- ▶ Each process begins life as a single thread called the *main thread*.
- ▶ At some point, the main thread creates a *peer thread*, and the two threads run concurrently.
- ▶ Thread context switch
 - Control passes to the peer thread via a context switch, because the thread executes a blocking I/O (or slow system call such as read, sleep or system's interval timer)

Posix threads (Pthreads)

- ▶ Posix threads (Pthreads) is a standard interface for manipulating threads from C programs.
- ▶ It was adopted in 1995 and is available on most Unix systems.

Creating Threads

- `int pthread_create(tid, attr, function pointer, arg)`
 - *tid*: Thread id (also accessible within thread by `pthread_self()`)
 - *attr*: Can be used to change the default attributes of the newly created thread (we will always call `pthread_create` with a NULL *attributes* argument)
 - *function pointer*: Expects function to execute
 - *arg*: Arguments passed as void pointer

```
#include <pthread.h>
typedef void *(func)(void *);

int pthread_create(pthread_t *tid, pthread_attr_t *attr,
                  func *f, void *arg);
```

Returns: 0 if OK, nonzero on error

Returning thread ID

- When `pthread_create` returns, argument `tid` contains the ID of the newly created thread. The new thread can determine its own thread ID by calling the `pthread_self` function

```
#include <pthread.h>  
  
pthread_t pthread_self(void);
```

Returns: thread ID of caller

Terminating Threads

- ❑ When function returns
- ❑ Explicit call to *pthread_exit()*
- ❑ When process terminates

```
#include <pthread.h>

void pthread_exit(void *thread_return);
```

Returns: 0 if OK, nonzero on error

Waiting for Thread Termination

- **thread_join(pthread_t tid, void **thread_return)**
 - The `pthread_join` function blocks until thread *tid* terminates, assigns the generic (`void *`) pointer returned by the thread
 - Can only wait for specific thread for which id is known

```
#include <pthread.h>
```

```
int pthread_join(pthread_t tid, void **thread_return);
```

Returns: 0 if OK, nonzero on error

Waiting for Thread Termination

- Detached thread: exit status not retained
 - `pthread_detach(tid)`
 - Detaches an existing thread

```
#include <pthread.h>  
  
int pthread_detach(pthread_t tid);
```

Returns: 0 if OK, nonzero on error

pthread_join() vs. pthread_detach()

❑ **pthread_join()**

- It is called from another thread (usually the thread that created it)
- Wait for a thread to terminate and obtain its return value

❑ **pthread_detach()**

- Can be called from either the thread itself or another thread
- Indicate that you do not want the thread's return value or wait for it to finish.

Example (Thread prints hello world): lecture_8_code/pthread_hello/helloworld.c

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

void *function_hello(void *vargp) /* Thread routine */
{
    printf("Hello, world!\n");
    printf("Thread id = %d \n", (int)pthread_self()); //print out the current thread ID
    return NULL;
}

int main(){
    pthread_t tid;
    //create a thread running thread function
    pthread_create(&tid, NULL, function_hello, NULL);
    printf("Waiting for the thread %d to terminate.\n", (int)tid);
    pthread_join(tid, NULL); //waits for the termination
    exit(0);
}
```

Terminal:
\$./main

Example (**Thread returns value**): lecture_8_code/pthread_return_values/main.c

```
#include <pthread.h>
#include <string.h>
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
void *function_integer(void *arg) {
    printf("Entering the function_integer() via thread: %d.\n", (int)pthread_self());
    printf("function_integer() returns.\n");
    return((void *)1);
}

void *function_string(void *arg) {
    printf("Entering the function_string() via thread: %d.\n", (int)pthread_self());
    printf("function_string() returns.\n");
    return((void *)"Hello");
}
```

Example (Thread returns value): lecture_8_code/pthread_return_values/main.c

```
int main()
{
    pthread_t tid1,tid2;
    void *tret;
    pthread_create(&tid1,NULL,function_integer,NULL); //create thread 1
    printf("Created thread: %d\n",(int)tid1);
    pthread_create(&tid2,NULL,function_string,NULL); //create thread 2
    printf("Created thread: %d\n",(int)tid2);

    pthread_join(tid1,&tret); //waiting for the thread 1 termination
    printf("Main process ID: %d. Thread 1's return value: %d\n",(int)getpid(),(int)tret);
    pthread_join(tid2,&tret); //waiting for the thread 2 termination
    printf("Main process ID: %d. Thread 2's return value: %s\n",(int)getpid(),(char *)tret);

    exit(0);
}
```

Terminal:
\$./main

Example (**pthread_exit** returns value): lecture_8_code(pthread_exit_return/main.c)

```
#include <stdio.h>
#include <pthread.h>
#include <stdlib.h>

void *function_integer(int s)
{
    printf("Entering function_integer() via thread: %d.\n", (int)pthread_self());
    printf("%d\n",s);
    pthread_exit("Terminate the current thread inside of the function_integer().");
    //terminate the current thread
}

void *function_string(char s[])
{
    printf("Entering function_string() via thread: %d.\n", (int)pthread_self());
    printf("%s\n",s);
    pthread_exit("Terminate the current thread inside of the function_string().");
    //terminate the current thread
}
```

Example (**pthread_exit** returns value): lecture_8_code/thread_exit_return/main.c

```
int main(void) {
    pthread_t id1,id2;
    void *a1,*a2;
    int i,ret1,ret2;
    int s1=1000;           //input arguments
    char s2[]="Hello world!"; //input arguments
    ret1=pthread_create(&id1,NULL, function_integer,(void *)s1);
    printf("Created thread: %d\n", (int)id1);
    ret2=pthread_create(&id2,NULL, function_string,(void *) s2);
    printf("Created thread: %d\n", (int)id2);

    if(ret1!=0){
        printf ("Create pthread1 error!\n");
        exit (1);
    }
    printf("Main thread is waiting for the thread %d to terminate.\n", (int)id1);
    pthread_join(id1,&a1);
```

Example (**pthread_exit** returns value): lecture_8_code(pthread_exit_return/main.c)

```
printf("return value from thread 1: %s\n",(char*)a1);

if(ret2!=0){
    printf ("Create pthread2 error!\n");
    exit (1);
}
printf("Main thread is waiting for the thread %d to terminate.\n", (int)id2);
pthread_join(id2,&a2);
printf("return value from thread 2: %s\n",(char*)a2);
return (0);
}
```

Terminal:
\$./main

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/server.c

```
#include <stdio.h>
#include <sys/socket.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <string.h>
#include <unistd.h>
#include <pthread.h>
#define PORT 8080

int main(int argc, char const *argv[])
{
    int server_fd, *new_socket;
    struct sockaddr_in address;
    int addrlen = sizeof(address);
    pthread_t tid;
```

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/server.c

```
// Creating socket file descriptor
if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0){
    perror("socket failed");
    exit(EXIT_FAILURE);
}
address.sin_family = AF_INET;
address.sin_addr.s_addr = htonl (INADDR_ANY);
address.sin_port = htons( PORT );
// Attaching socket to the port 8080
if (bind(server_fd, (struct sockaddr *)&address, sizeof(address))<0) {
    perror("bind failed");
    exit(EXIT_FAILURE);
}
if (listen(server_fd, 1024) < 0) { //listen to the socket
    perror("listen");
    exit(EXIT_FAILURE);
}
```

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/server.c

```
while(1){  
    printf("Waiting for a client to connect...\n");  
    new_socket = malloc(sizeof(int));  
    if ((*new_socket = accept(server_fd, (struct sockaddr *)&address,  
        (socklen_t*)&addrlen))<0)  
    {  
        perror("accept");  
        exit(EXIT_FAILURE);  
    }  
  
    pthread_create(&tid, NULL, thread_handle, new_socket);  
    pthread_detach(tid);  
}  
  
return 0;  
}
```

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/server.c

```
void *thread_handle(void *vargp){  
    printf("Entering the thread: %d\n", (int)pthread_self());  
    char buffer[1024] = {0};  
    char *response = "Response from server";  
    int connfd = *((int *)vargp);  
    int valread = read( connfd , buffer, 1024); //receive the message from client  
    printf("%s\n",buffer );  
    write(connfd, response, strlen(response)); //send response from server  
    printf("Response msg sent from the thread: %d.\n", (int)pthread_self());  
    free(vargp);  
    close(connfd);  
    return NULL;  
}
```

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/client.c

```
#include <stdio.h>
#include <sys/socket.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#define PORT 8080

int main(int argc, char const *argv[])
{
    struct sockaddr_in address;
    int sock1, sock2, valread;
    struct sockaddr_in serv_addr;
    char *hello1 = "Hello from client 1";
    char *hello2 = "Hello from client 2";
    char buffer[1024] = {0};
```

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/client.c

```
if ((sock1 = socket(AF_INET, SOCK_STREAM, 0)) < 0)
{
    printf("\n Socket creation error \n");
    return -1;
}
if ((sock2 = socket(AF_INET, SOCK_STREAM, 0)) < 0)
{
    printf("\n Socket creation error \n");
    return -1;
}
// initialize serv_addr
memset(&serv_addr, '0', sizeof(serv_addr));

serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(PORT);
```

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/client.c

```
if/inet_pton(AF_INET, "127.0.0.1", &serv_addr.sin_addr)<=0) {  
    printf("\nInvalid address/ Address not supported \n");  
    return -1;  
}  
if (connect(sock1, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0) {  
    printf("\nConnection Failed \n");  
    return -1;  
}  
// send messages via two sockets: sock1  
write(sock1 , hello1 , strlen(hello1));           //send msg via sock1  
printf("Hello message sent from client 1\n");  
valread = read(sock1, buffer1, 1024);             //receive msg from server  
printf("%s\n", buffer1);
```

Example (**Concurrent server via pthread**): lecture_8_code/concurrent_server_thread/client.c

```
if (connect(sock2, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0){  
    printf("\nConnection Failed \n");  
    return -1;  
}  
// send messages via two sockets: sock2  
write(sock2 , hello2 , strlen(hello2));           //send msg via sock2  
printf("Hello message sent from client 2\n");  
valread = read(sock2, buffer2, 1024);             //receive msg from server  
printf("%s\n", buffer2);  
return 0;  
}
```

The server can respond to multiple clients' requests concurrently through creating threads

Terminal 1
\$./server

Terminal 2
\$./client

Terminal 3
\$./client

Race Conditions

- Unpredictable interleaved execution can lead to race conditions when accessing shared resources
 - Between processes (e.g., file access)
 - Between threads (global, static variables)
- Thread libraries provide solutions for threads

Writing Thread-Safe Functions

- ▶ Problem: Function refers to global or static variables



Solutions?

- ▶ Local variables are safe, because each Thread has own stack

Mutex

- ▶ Mutual exclusion lock – only one thread can acquire the lock at any given time
 - Other threads that try to lock will block
- ▶ `int pthread_mutex_lock(pthread_mutex_t *mutex)`
- ▶ `int pthread_mutex_unlock(pthread_mutex_t *mutex)`

The macro `PTHREAD_MUTEX_INITIALIZER` can be used to initialize mutex that are statically allocated.

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

Example (**mutex example**): lecture_8_code/mutex/main.c

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#define NLOOP 500
/*
Two threads: each thread will increase the global counter NLOOP times, and
print it out with the thread ID.
*/
int counter; /* this is incremented by the threads */
//mutex initialization
pthread_mutex_t counter_mutex = PTHREAD_MUTEX_INITIALIZER;
```

Example (**mutex example**): lecture_8_code/mutex/main.c

```
void *doit(void *vptr) {
    printf("Entering the thread %d\n", (int)pthread_self());
    int i, val;
    for (i = 0; i < NLOOP; i++) {
        pthread_mutex_lock(&counter_mutex);
        val = counter;
        printf("%d: %d\n", (int)pthread_self(), val + 1);
        counter = val + 1;
        pthread_mutex_unlock(&counter_mutex);
    }
    return(NULL);
}
```

Example (**mutex example**): lecture_8_code/mutex/main.c

```
int main(int argc, char **argv) {
    pthread_t tidA, tidB;
    pthread_create(&tidA, NULL, doit, NULL); //create thread 1
    printf("Created thread %d\n", (int)tidA);
    pthread_create(&tidB, NULL, doit, NULL); //create thread 2
    printf("Created thread %d\n", (int)tidB);

    /*wait for both threads to terminate */
    pthread_join(tidA, NULL);
    printf("Thread %d is terminated.\n", (int)tidA);
    pthread_join(tidB, NULL);
    printf("Thread %d is terminated.\n", (int)tidB);
    exit(0);
}
```

Terminal
\$./main

Example (**non_mutex** example): lecture_8_code/non_mutex/main.c

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#define NLOOP 500
/*
Two threads: each thread will increase the global counter NLOOP times, and
print it out with the thread ID.
*/
int counter; /* this is incremented by the threads */
//mutex initialization
pthread_mutex_t counter_mutex = PTHREAD_MUTEX_INITIALIZER;
```

Example (**non_mutex** example): lecture_8_code/non_mutex/main.c

```
void *doit(void *vptr) {
    printf("Entering the thread %d\n", (int)pthread_self())
    int i, val;
    for (i = 0; i < NLOOP; i++) {
        //pthread_mutex_lock(&counter_mutex);
        val = counter;
        printf("%d: %d\n", (int)pthread_self(), val + 1);
        counter = val + 1;
        //pthread_mutex_unlock(&counter_mutex);
    }
    return(NULL);
}
```

Example (**non_mutex** example): lecture_8_code/non_mutex/main.c

```
int main(int argc, char **argv) {
    pthread_t tidA, tidB;
    pthread_create(&tidA, NULL, doit, NULL); //create thread 1
    printf("Created thread %d\n", (int)tidA);
    pthread_create(&tidB, NULL, doit, NULL); //create thread 2
    printf("Created thread %d\n", (int)tidB);

    /*wait for both threads to terminate */
    pthread_join(tidA, NULL);
    printf("Thread %d is terminated.\n", (int)tidA);
    pthread_join(tidB, NULL);
    printf("Thread %d is terminated.\n", (int)tidB);
    exit(0);
}
```

Terminal
\$./main

Homework Readings

- ▶ CSAPP: Computer Systems – A Programmer’s Perspective
 - Chapter 12.1, 12.3, 12.5
- ▶ ALP: Advanced Linux Programming
 - Chapter 8.3

Assignment 3

Concurrent Web Proxy

3 weeks

Due on 4/22

Problem

- ▶ You are required to upgrade your sequential Basic Web Proxy of Assignment 2 so that it uses threads or processes to deal with multiple clients concurrently.
- ▶ You will be required to implement **two versions** of the concurrent web proxy:
 - A threaded version and a multi-process version

Implementation

- ▶ Real proxies do not process requests sequentially. They deal with multiple requests concurrently. Once you have a working sequential logging proxy (Assignment 2), you should alter it to handle multiple requests concurrently. Following are two approaches to achieve concurrency:
- ▶ **Process based** using *file locking* for synchronization: Create a new process to deal with each new connection request that arrives.
- ▶ **Thread based** using *mutexes* for synchronization: Create a new thread to deal with each new connection request that arrives.